

EC504 ALGORITHMS AND DATA STRUCTURES
FALL 2020 MONDAY & WEDNESDAY
2:30 PM - 4:15 PM

Prof: David Castañón, dac@bu.edu

GTF: Mert Toslali, toslali@bu.edu

Haoyang Wang: haoyangw@bu.edu

Christopher Liao: cliao25@bu.edu

Project: Due Friday

Report: e-mail to both:

dec@bu.edu, tosh@bu.edu

Also: tell us where the
software is!

Exam 2: Th 5/6 3-5

F 5/7 9-11

E-mail if: you will take it
on F 5/7.

elif: I will assign you to
Th 3-5.

else: ☹ grade...

Please fill a Course eval...(today)

→ Important ABET visit
next year...

Graph Terminology

- Directed, Undirected
- Paths, simple paths, cycles, simple cycles
- Connected graphs
- Trees, forests
- Euler paths, existence of Euler cycles
- Degree of a vertex, neighbors of vertices
- Connectivity in directed paths
- Graph representations
 - Adjacency list, forward star, adjacency matrix

Graph Algorithms

- Graph traversals
 - BFS, DFS, how they are used for solving
- Minimum Spanning Trees
 - Prim's, Kruskal's, Boruvka's
- Single source shortest path algorithms
 - Dijkstra's, Bellman-Ford •
- All pairs shortest paths •
 - Floyd-Warshall, Johnson's
- Single source, single destination
 - A* search •

Breadth-First Search (uses Queue)

1. Mark all vertices as unvisited, parents as NULL, depth as -1
 2. Choose any unvisited vertex, mark it as **visited** and enqueue it onto queue
 3. While the queue is not empty:
 - Dequeue top vertex v from the queue. Do work to be done on that vertex
 - If $\text{parent}[v] == \text{NULL}$, set depth to 0; otherwise, set depth to $\text{depth}[\text{parent}[v]] + 1$
 - For each vertex adjacent to v (e.g. in out list) that has not been visited: Mark it visited, mark its parent as v , and enqueue it
 - Mark v as done
 4. If there are unvisited vertices, choose any unvisited vertex, mark it as **visited**, enqueue it and repeat step 3
- This can handle graphs that are not connected
 - Marking as visited avoids cycles
 - Complexity: $O(\#V + \#E)$, reduces to $O(\#E)$ if strongly connected
 - Size of queue is $O(\#V)$

Depth-First Search

Stacks




Recursive implementation:

1. Mark all vertices as unvisited; mark all parents as NULL
2. While there are vertices marked as unvisited:
 - Select unvisited vertex v , mark as visited:
 - Do DFS(vertex)


DFS(vertex):

- For neighbors of vertex
 - If neighbor is unvisited, mark as visited and do DFS(neighbor)
- This can handle graphs that are not strongly connected
 - Marking as visited avoids cycles
 - Complexity: $O(\#V + \#E)$, reduces to $O(\#E)$ if strongly connected
 - Size of queue is $O(\#V)$

Example: BFS



Queue




Example

Performing a recursive depth-first traversal:


- Insert A: Visited: A, B
Stack: A, B

Examine B: Visited A, B, C,
Stack A, B, C



DFS for Topological Sort

- Alternative algorithm: **recursive** DFS ($O(\#E)$)
 - Order in which vertices are completed is reverse order of a topological sort!



Stack: 1 2 4	Completed: 4
Stack: 1 2 3	Completed: 3
Stack: 1 2 5	Completed: 5
Stack: 1 2	Completed: 2
Stack: 1	Completed: 1
Stack: 6	Completed: 6
Reverse Order: 6, 1, 2, 5, 3, 4	

DFS for Biconnectivity

- Connected **undirected** graph is **biconnected** if there are no nodes whose removal disconnects the graph
- Nodes whose removal disconnect the graph are known as **articulation points**
- DFS can be used to find articulation points:
 - Algorithm: Number nodes in Depth First Search order, in the order in which they are inserted into the execution stack of the recursive Depth-First Search. This creates a spanning tree in the graph. Call this number $NUM(n)$ for node n


DFS for Strongly-Connected Components

○ Kosaraju's Algorithm

- Perform DFS on graph $G = (V, E)$,
 - Number vertices according to their finishing time in DFS of G
- Perform DFS on $G_r = (V, E_r)$, where E_r are reverse of edges in E , selecting nodes to start in the stack, in decreasing order of finishing time in previous DFS
- Strongly connected components = reachable trees obtained in last DFS


Weighted Graphs

- A weighted graph $G = (V, E)$ is a graph along with a weight function $w : E \rightarrow \mathfrak{R}$
- Weighted graphs can be directed or undirected



Spanning Trees


- ▶ A **spanning tree** of an undirected graph is
 - ▶ edge subset forming a tree that spans every vertex, has $\#V - 1$ edges
- ▶ A **minimum spanning tree** (MST) of an undirected weighted graph (V, E) with weights $w(\cdot)$ is a spanning tree with the smallest sum of the weights of its edges



Kruskal's Algorithm

- Sort edges by weight in ascending order
- Start with empty set T (note: it is promising)
- For each edge e in sorted list
 - If adding edge e to T does not create cycle in $(V, T \cup e)$
 - ...add it to MST: $T = T \cup \{e\}$
 - Claim: T is now promising set with one more edge
- Stop when you have $\#V - 1$ edges in T

Example



Yes

Done

edges = [(C,E), (D,F), (B,C), (E,F), (B,D), (A,B), (A,D), (B,E), (B,F)]


Kruskal Runtime

- $O(|V|)$ for iterating through vertices
- $O(|E|\log|E|)$ for sorting edges
- $O(|E|\times 1)$ for iterating through edges and merging clouds with path compression
- $O(|V|+|E|\log|E|+|E|\times 1)$
 - $= O(|V|+|E|\log|E|)$
- $O(|V|+|E|\log|E|)$ ✓
 - Better than simple $O(|V|^3)$ without disjoint sets

Prim-Jarnik

- Traverse $G = (V, E)$ starting at any node
 - Maintain priority queue of nodes (e.g. binary heap, Fibonacci heap)
 - set priority to weight of the cheapest edge that connects them to MST
- Un-added nodes start with priority ∞
- At each step *edge*
 - Add the ~~node~~ with lowest cost to MST
 - Update ("relax") neighbors as necessary
- Stop when all nodes added to MST

Example



$PQ = [(0, A), (\infty, B), (\infty, C), (\infty, D), (\infty, E), (\infty, F)]$

Runtime

- Initializing nodes with distance and previous pointers is $O(|V|)$; putting nodes in PQ is $O(|V|)$
- While loop runs $|V|$ times
 - removing vertex from PQ is $O(\log|V|)$
 - So $O(|V|\log|V|)$
- For loop (in while loop) runs $|E|$ times in total
 - Determining whether v' is in PQ: $O(1)$ if we build index into PQ
 - Decreasing vertex's key in the PQ is $\log|V|$ (binary heap), or amortized to $O(1)$ if we use Fibonacci or rank-pairing heaps
 - So $O(|E|)$ in complex data or $O(|E| \log|V|)$
- Overall runtime
 - $O(|V| + |V|\log|V| + |E|) = O(|E| + |V|\log|V|)$ ✓


$|E|/|V|$

Borůvka's Algorithm

- Earliest MST algorithm: 1926. Application: design of power grid
- For **every** connected set in a forest,
 - Select smallest weight edge that leaves connected set
 - Add it to the MST
 - But don't add it twice if same edge selected by two connected sets
- In principle, merges at least half of the trees at each time: $O(\log(n))$ iterations
 - Easy to parallelize ✓
- Each pass is $O(\#E)$

Example

- Start with every vertex in a separate connected set, partial MST empty



MST Algorithms are Greedy Algorithms

- ▶ Greedy algorithm: an algorithm that builds a solution adding an element at a time that is the locally optimal choice at that time
 - ▶ Uses a simple rule e.g. MST adds edge with minimum weight across a cut
 - ▶ No backtracking (!!!)
 - ▶ Greedy algorithms are not always optimal . . .
 - ▶ / Special classes of problems can be solved to optimality by greedy algorithms

Other Problems with Greedy Algorithm Solutions


▸ Scheduling

- Interval scheduling: one processor, jobs with start and end times, maximize number of jobs done: ~~schedule by earliest finish time~~
- Scheduling to minimize lateness; jobs with start times, deadlines. Minimize maximum lateness → Earliest deadline first scheduling
- Single server, N jobs with different processing times, all available to start right away: to minimize sum of finishing times over jobs → smallest processing time first
- Single server, N jobs, unit processing times, hard deadlines $d(j)$, value of scheduling $V(j)$: Try to add to feasible set in order of decreasing value, verify feasible using just-in-time schedule
- Fractional Knapsack: schedule in order of decreasing value/size ratio.

yes


Single Source Shortest Paths (SSSP)

- ▶ Given a graph and a source vertex
 - ▶ find the shortest paths to all other vertices
 - ▶ results in a shortest-path tree
 - ▶ Single directed path to every other vertex




Important Property of Shortest Paths

- ✦ **Lemma:** The shortest path from vertex s to a vertex t is composed of shortest paths to and from any intermediate vertices
- ▶ Bellman's Principle of Optimality
- ▶ Leads to dynamic programming




Dijkstra's Algorithm

- ▶ **Greedy** Algorithm: Assumes all edges have **nonnegative** weights
 - ▶ Maintain a set of **explored** nodes S for which algorithm has determined $D[u]$ = length of a shortest s to u path
 - ▶ Initialize $S = \{s\}$, $D[s] = 0$; $D[v] = \text{infinity}$
 - ▶ Choose unexplored node $v \notin S$ which minimizes $D[u] + w(u, v)$, $u \in S$
Set $D[v] = \min_{\{(u,v) \in E: u \in S\}} D[u] + w(u, v)$, and add v to S
Set $\text{pred}[v] = \text{vertex } u \text{ in } S \text{ that achieves } d[v]$
 - ▶ Repeat until all vertices are explored, so $S = V$
 - ▶ Path to any vertex can be found by using $\text{pred}[]$ labels
- ▶ Complexity $O(|E| + |V| \log(|V|))$ with Fibonacci or Rank-pairing heaps



Another Example



	A	B	C	D	E
D[]	0	∞	∞	∞	∞
predecessor[]					

A
B
D
E

4
3


2
5

Bellman-Ford Algorithm

- Algorithm converges to shortest paths in graphs with negative distances, provided no negative cycles exist in graph: $O(|E||V|)$ worst case with very simple data structures
- Use a queue of vertices where distances have changed
 - Initialize $D[s] = 0$; $D[v] = \infty, v \neq s$; $\text{pred}[v] = \text{null}$
 - Insert s into queue Q; mark $\text{inqueue}[s] = \text{true}$, mark $\text{inqueue}[v] = \text{false}, v \neq s$
 - While Q is not empty:
 - Select u out of queue, mark $\text{inqueue}[u] = \text{false}$
 - For each edge (u,v) in E:
 - If $D[v] > D[u] + w(u, v)$:
 - Set $D[v] = D[u] + w(u, v)$, set $\text{pred}[v] = u$
 - if $\text{inqueue}[v] = \text{false}$, add v to Q, mark $\text{inqueue}[v] = \text{true}$

Another Example

BF



	A	B	C	D	E
D[]	0	∞	∞	∞	∞
pred[]		4	2		

A
B
C
D
E

B

3..

28

6 7

5 6

~~A~~
~~B~~
~~C~~
~~D~~
~~E~~

All Pairs Shortest Paths

Input: Directed graph $G = (V, E)$, where $V = \{1, 2, \dots, n\}$, with edge-weight function $w : E \rightarrow \mathbb{R}$

- Weights may be negative, but no negative cycles ✓

Output: $n \times n$ matrix of shortest-path lengths $D(i, j)$ for all $i, j \in V$ (and routes)

Floyd-Warshall: initially $D(i, j) = \begin{cases} w(i, j) & \text{if } (i, j) \in E \\ \infty & \text{if } (i, j) \notin E \end{cases}$

- For k in V : For i in V : For j in V :

 If $D(i, j) > D(i, k) + D(k, j)$:

 set $D(i, j) = D(i, k) + D(k, j)$ ↩

- Complexity: $O(|V|^3)$



All Pairs Shortest Paths: Alternative


Alternative: Run Dijkstra's algorithm from every possible starting edge

Problem: negative edges

Can fix with Johnson's algorithm: preprocess with Bellman-Ford and rescale edge distances using weights

- **Complexity** $O(|V|(|E| + |V| \log(|V|)))$ with fancy heaps

Reweighting



Shortest Path from s to a single t

- A* search
 - Will search in direction of t for shortest path; a modification of Dijkstra
- Need heuristic function to estimate distance remaining to t
 - admissible: for each vertex v: $h(v) \leq D(v,t)$, where $D(v,t)$ is best distance
 - consistent: For each (u,v) in E: $h(u) \leq w(u,v) + h(v)$
- Admissible, consistent heuristic leads to optimality

Maximum Flow in Flow Networks

- A feasible st-flow (flow) $f: E \rightarrow \mathbb{R}^+$ is a function that satisfies


$0 \leq f(u, v) \leq c(u, v)$ for all edges $(u, v) \in E$ (capacity satisfied)

For every $v \in V - \{s, t\}$, flow is conserved: $\sum_{(u,v) \in E} f(u, v) = \sum_{(v,w) \in E} f(v, w)$

- Value of flow: net flow out of s


$$val(f) = \sum_{(s,v) \in E} f(s, v) - \sum_{(v,s) \in E} f(v, s)$$

- Maximum flow problem: find flow f that maximizes $val(f)$
- Value of max-flow = value of min-cut
- Any s-t cut is an upper bound to the max flow problem




Residual Network

- Current solution



- Residual network has capacity




BFS. \rightarrow Polynomial

Ford-Fulkerson Algorithm

- Start with flow $f(u,v) = 0, (u, v) \in E$. Form the residual network $G_f = G$
- While there exists an $s \rightarrow t$ path P in the residual network
 - Compute residual capacity δ on P , and augment flow f using flow δ on path P
 - Update residual network using new flow f , as G_f
- When no path can be found, return flow f
- Complexity $O(|V||E|C)$, where C is largest capacity: pseudopolynomial
- Variation: Edmonds-Karp \rightarrow Find minimum hop augmenting paths, guaranteed polynomial complexity $O(|V||E|^2)$

Example

- Numbers are capacities, $s = 1$, $t = 8$



- BFS: augment $1 \rightarrow 2 \rightarrow 7 \rightarrow 8$, capacity 6

Preflow-Push Algorithms ('88)

- A **preflow** is a function $x : E \rightarrow \mathbb{R}^+$, where $0 \leq x(u, v) \leq c(u, v)$ and

$$\text{e}(v) = \sum_{(u,v) \in E} x(u, v) - \sum_{(v,w) \in E} x(v, w) \geq 0, \text{ for } v \in V - \{s, t\}$$

- $\text{e}(v)$ is the excess at vertex v , required to be non-negative
- Let G_x be residual network for a preflow $x()$. Distance labels $d()$ are **valid** for G_x if $d(t) = 0$ and $d(v) \leq d(u) + 1$ for each $(u, v) \in G_x$
- Let $r(u, v)$ be the capacity of edge (u, v) in residual network G_x . An edge (u, v) is **admissible** if $r(u, v) > 0$ and $d(u) = d(v) + 1$

$$d(u) \leq d(v) \rightarrow !$$

Goldberg-Tarjan Preflow Push Algorithm

- **Initialize:**

- $G_x = G; x(u, v) = 0, (u, v) \in E$
- Using BFS reverse from t , compute distance $d(v)$ for every vertex v
- For every $(s, v) \in E$, set $x(s, v) = c(s, v)$; set $e(v) = c(s, v)$; set $d(s) = |V|$
- Update G_x , with residual capacities $r(u, v), (u, v) \in E_x$
- While there is an active node in G_x , select active vertex v and push/relabel(v):
 - If there is admissible edge (v, w) : $x(v, w) := x(v, w) + \min(e(v), r(v, w))$
 - Otherwise increase $d(v)$: $d(v) = \min\{d(w) + 1 : (v, w) \in E_r\}$
- Once there are no active nodes, send all excess flow back to s

✓ Dynamic Programming (DP)

1 New Path
DP...

- A general approach for breaking solutions of large problems into sequence of solutions of smaller problems
 - Used in shortest path algorithms (BF, FW, A*)
- Weighted interval scheduling, maximum subarray sum, rod cutting
 - Examples of how to use DP for new problems
- Integer Knapsack: Solvable by DP
 - pseudo-polynomial $O(nC)$, C is knapsack size
- Sequence Alignment
 - Solvable by DP, polynomial $O(mn)$

j \ k	0	1	2	3	4	5	6	7	8	9	10	11
1	0	1	1	1	1	1	1	1	1	1	1	1
2	0	1	6	7	7	7	7	7	7	7	7	7
3	0	1	6	7	7	18	19	24	25	25	25	25
4	0	1	6	7	7	7	22	23	28	29	29	40
5	0	1	6	7	7	7	22	28	29	34	35	40

	x_1	x_2	x_3	x_4	x_5	x_6
	C	T	A	C	C	G
	-	T	A	C	A	T
	y_1	y_2	y_3	y_4	y_5	y_6

an alignment of CTACCG and TACATG

$M = \{x_2 - y_1, x_3 - y_2, x_4 - y_3, x_5 - y_4, x_6 - y_5\}$


$$\text{Cost}(M) = \sum_{i,j \text{ matched}} s(x_i, y_j) + \sum_{i: x_i \text{ unmatched}} \delta + \sum_{j: y_j \text{ unmatched}} \delta$$

Assignment Problems


$$\min_{\{x_{ij} \in \{0,1\}\}} \sum_{(i,j) \in E} C_{ij} x_{ij} \text{ subject to constraints}$$

$$\sum_{i:(i,j) \in E} x_{ij} = 1, j \in 1, \dots, n; \quad \sum_{j:(i,j) \in E} x_{ij} = 1, i \in 1, \dots, n$$

Graph can be sparse, edges have capacity 1



Equivalent network flow representation




Successive Shortest Path Algorithm

- Define prices $q_i, i = 1, \dots, n$ for persons and prices $p_j, j = 1, \dots, n$ for objects
- Define reduced costs of edges with those prices as $c_{p,q}^r(i, j') = C_{ij'} + q_i - p_{j'}$
- Initially, set $x_{ij'} = 0, (i, j') \in E$, matching $M' = \emptyset$, set prices $p_{j'} = \min_{(i, j') \in E} C_{ij'}$, $q_i = 0, i, j' \in 1, \dots, n$
- Construct the residual network (V, E^r) given matching M' and the prices $\{p, q\}$
 - Cost of arcs $(i, j) \in E$: $c_{p,q}^r(i, j)$; cost of reverse arcs (j, i) , where $(i, j) \in E$: $-c_{p,q}^r(i, j)$
- Find shortest augmenting path P from s to t; compute the shortest distances $d(i), d(j')$ to vertices $i=1, \dots, n, j' = 1, \dots, n$
- Raise prices $q_i := q_i + d(i), i = 1, \dots, n$; $p_{j'} := p_{j'} + d(j'), j' = 1, \dots, n$
- Modify assignments on augmenting path P by one unit ✓
- Repeat above iteration n times until complete matching is found
- Complexity $O(|V||E| + |V|^2 \log(|V|))$


Do once or twice

Data Structures for Multidimensional Search

- PR Quadtrees are tries ✓
 - Children of a node: four quadrants of partition of a region
 - If a leaf has more than one point, it splits into 4 subregions
 - Insert, delete, search



- k-d trees
 - binary search tree where branching decisions are made based on different coordinates at each level
 - Batch construction using medians result in balanced trees
 - One-by-one insertion can be very unbalanced
 - Insertion, deletion, search ✓



R-trees

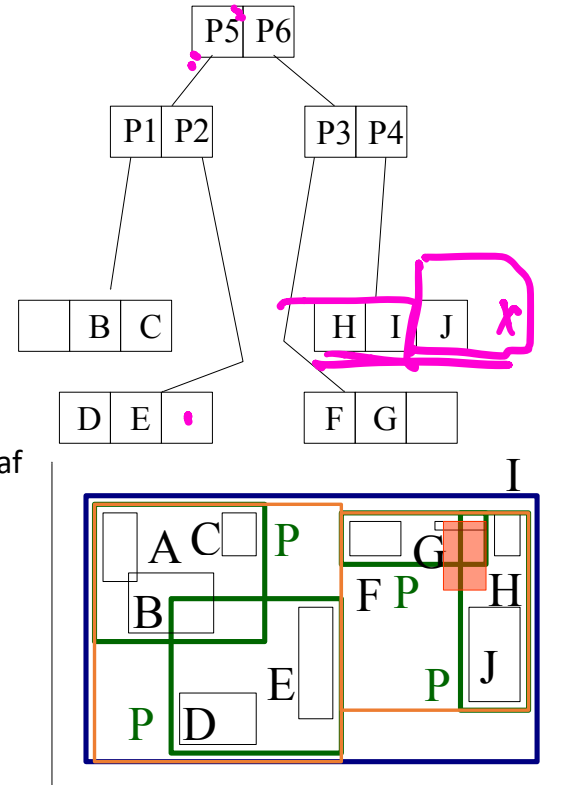
~ B+ trees

- Storage for regions

- Keys: n-dimensional rectangles, (2 points)
- All leaf nodes appear on the same level
- Every node contains between m and M entries
 - $m \leq M/2$ is the minimum entries per node
- Root node has at least 2 entries (children)

- Insert

- Insert into rectangle that increases the least by adding
- Increase measure by perimeter or area — descend to leaf
- If node saturates (M+1) entries, must split
 - Linear or ~~Quadratic~~ criteria to pick seeds
 - Add to seeds by smallest increase in area
 - Guarantee minimum of m in each split




Computation Complexity

- **Decision problems** have a yes or no solution. Abstract decision problem is a function which maps problem instances I into $\{\text{yes}, \text{no}\}$
- Deterministic, non-deterministic Turing machines ...
- Class **P**: there is a algorithm solving the problem with a running time on a deterministic machine that is polynomial in the input size
- Class **NP** (non-deterministic polynomial): given a candidate solution, there is a polynomial-complexity algorithm to verify whether the answer is yes or no for that solution
- Problem A is polynomially reducible to problem B if there exists an algorithm for solving problem A in polynomial time if we could solve arbitrary instances of problem B at unit cost ($A \leq_P B$)
 - If $A \in \mathbf{P}$ and $B \leq_P A$, then $B \in \mathbf{P}$
- Problem A is NP-complete, if $A \in \mathbf{NP}$ and every problem $B \in \mathbf{NP}$ can be polynomially reduced to A. That is, $B \leq_P A$

$P \stackrel{?}{=} NP$

Other Complexity Concepts

- **NP hard problem:** there is an NP-complete problem Y , such that Y is reducible to X in polynomial time (but X may not be in NP)
- **Pseudopolynomial** complexity: If K is the size of the largest number, and n is the size of the input, then the worst case complexity is polynomial in n and K (e.g. Integer knapsack)
- **Strongly Polynomial** complexity: worst case complexity is polynomial in input size, independent of largest value of number in input
- **Strongly NP-complete problems:** If one restricts the size of the largest number in the problem to K , where K is a polynomial in the input size n , then the problem is still NP-complete
 - e.g. Clique



No PTAS
Yes for Pseudopoly. —

Approximate Algorithms

- Objective: Find approximate algorithms for NP-hard problems with performance guarantees
 - Solution of approximate algorithm is within a factor of optimal solution ✓
- Integer Knapsack: Greedy achieves at least 50% of optimal value ✓
- TSP: MST heuristics can generate tour that is no longer than twice the distance D^* of the optimal TSP tour (can improve to 1.5) ✓
- For pseudopolynomial complexity problems, can usually approximate within epsilon in time that grows as $1/\epsilon$ using rounding

PTAS

X

$$E \log E = E \log V$$

$$U \subseteq E$$

$$V \subseteq E$$

$$E = \underline{U}^2$$

