

EC504 ALGORITHMS AND DATA STRUCTURES
FALL 2020 MONDAY & WEDNESDAY
2:30 PM - 4:15 PM

Prof: David Castañón, dac@bu.edu

GTF: Mert Toslali, toslali@bu.edu

Haoyang Wang: haoyangw@bu.edu

Christopher Liao: cliao25@bu.edu

Representations of graphs

- Adjacency matrix:


- $A = [a_{ij}]$, where

$$a_{ij} = \begin{cases} 1 & \text{if } (i, j) \in E \\ 0 & \text{if } (i, j) \notin E \end{cases}$$

- Symmetric for undirected graphs

$$A = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{bmatrix} \begin{matrix} \cdots 1 \\ \cdots 2 \\ \cdots 3 \\ \cdots 4 \\ \vdots \\ \vdots \end{matrix}$$

1 2 3 4



Adjacency List

- Each vertex has list of neighbor nodes

- For directed graphs, out-list

- Used in MATLAB for sparse matrices


- Can be implemented as arrays

- Forward-Star

- or linked lists


- $|V| + |E|$ for directed graphs

- Must store 2 edges for undirected



```
1 • → 2 → 4
2 •
3 • → 5
4 • → 2 → 5
5 • → 2 → 3 → 8
6 • → 9
7 • → 9
8 • → 4
9 •
```

V vertices
E edges





1. Edge list: {a,b}, {b,c}, {c,d}, {d,e}, {e,a}, {e,b}

2. Adjacency list

Vertex	Adjacencies
a	b, e
b	a, c, e
c	b, d
d	c, e
e	a, b, d

Linked List Form





Forward Star Representation

Vertex Array First[I]

i = 1

1

i = 2

3

i = 3

4

i = 4

4

i = 5

6

i = 6

7

Edges[j]

j = 1

2

j = 2

5

j = 3

3

j = 4

3

j = 5

5


j = 6

2

j = 7

-1

Just the adjacency list put end to end in the arc array!



3. Adjacency matrix:

A=

	a	b	c	d	e
a	0	1	0	0	1
b	1	0	1	0	1
c	0	1	0	1	0
d	0	0	1	0	1
e	1	1	0	1	0

4. Incidence matrix:

Matrix of vertex rows, edges columns

M=

Directed: -1 in start of edge, 1 in end

Undirected: 1 in both

	1	2	3	4	5	6
a	1	1	0	0	0	0
b	0	1	1	0	1	0
c	0	0	0	0	1	1
d	0	0	0	1	0	1
e	1	0	1	1	0	0




Figure 22.1 Two representations of an undirected graph. (a) An undirected graph G having five vertices and seven edges. (b) An adjacency-list representation of G . (c) The adjacency-matrix representation of G .




Figure 22.2 Two representations of a directed graph. (a) A directed graph G having six vertices and eight edges. (b) An adjacency-list representation of G . (c) The adjacency-matrix representation of G .

Graph Traversals

- Traversals of graphs are also called *searches*
- We can use either breadth-first or depth-first traversals
 - Breadth-first requires a queue
 - Depth-first requires a stack
- In each case, we will have to track which vertices have been visited requiring $\Theta(|V|)$ memory
- The time complexity cannot be better than and should not be worse than $\Theta(|V| + |E|)$
 - Connected graphs simplify this to $\Theta(|E|)$
 - Worst case: $\Theta(|V|^2)$

Breadth-First Search

1. Mark all vertices as unvisited, parents as NULL, depth as -1
 2. Choose any unvisited vertex, mark it as **visited** and enqueue it onto queue
 3. While the queue is not empty:
 - Dequeue top vertex v from the queue. Do work to be done on that vertex
 - If $\text{parent}[v] == \text{NULL}$, set depth to 0; otherwise, set depth to $\text{depth}[\text{parent}[v]] + 1$
 - For each vertex adjacent to v (e.g. in out list) that has not been visited: Mark it visited, mark its parent as v , and enqueue it
 - Mark v as done
 4. If there are unvisited vertices, choose any unvisited vertex, mark it as **visited**, enqueue it and repeat step 3
- This can handle graphs that are not connected
 - Marking as visited avoids cycles
 - Complexity: $O(\#V + \#E)$, reduces to $O(\#E)$ if strongly connected
 - Size of queue is $O(\#V)$

Depth-First Search

Recursive implementation:

1. Mark all vertices as unvisited; mark all parents as NULL
2. While there are vertices marked as unvisited:
 - Select unvisited vertex v , mark as visited:
 - Do DFS(vertex)

DFS(vertex):

- For neighbors of vertex
 - If neighbor is unvisited, mark as visited and do DFS(neighbor)
- This can handle graphs that are not strongly connected
 - Marking as visited avoids cycles
 - Complexity: $O(\#V + \#E)$, reduces to $O(\#E)$ if strongly connected
 - Size of queue is $O(\#V)$

Recursive depth-first traversal

A recursive implementation uses the call stack for memory:

```
void Graph::depth_first_traversal( Vertex *first ) const {
    std::unordered_map<Vertex *, int> hash;
    hash.insert( first );

    first->depth_first_traversal( hash );
}

void Vertex::depth_first_traversal( unordered_map<Vertex *, int> &hash ) const {
    // Perform an operation on this

    for ( Vertex *v : adjacent_vertices() ) {
        if ( !hash.member( v ) ) {
            hash.insert( v );
            v->depth_first_traversal( hash );
        }
    }
}
```

Depth-First Search

Iterative implementation: (Similar to Recursive)

1. Mark all vertices as unvisited; mark all parents as NULL
2. While there are vertices marked as unvisited:
 - Choose any unvisited vertex, mark it as **visited** and push it onto stack
 - While stack is not empty:
 - Let v be the vertex that is on top of stack.
 - If v has no unvisited neighbors, pop v , work on it and mark it done
 - Else select an unvisited neighbor of v , mark neighbor as visited, set $\text{parent}[\text{neighbor}] = v$, and push neighbor onto stack
- This can handle graphs that are not strongly connected
 - Marking as visited avoids cycles
 - Complexity: $O(\#V + \#E)$, reduces to $O(\#E)$ if strongly connected
 - Size of queue is $O(\#V)$

Depth-First Search

Alternative Iterative implementation: (Different from Recursive)

1. Mark all vertices as unvisited; mark all parents as NULL
2. While there are vertices marked as unvisited:
 - Choose any unvisited vertex v , mark it as **visited** and push it onto stack
 - While stack is not empty:
 - Pop v be the vertex that is on top of stack. Work on v . For any unvisited neighbors of v in its out list, mark neighbor as visited, set $\text{parent}[\text{neighbor}] = v$, and push it onto stack. Mark v as done
- This can handle graphs that are not strongly connected
 - Marking as visited avoids cycles
 - Complexity: $O(\#V + \#E)$, reduces to $O(\#E)$ if strongly connected
 - Size of queue is $O(\#V)$

Iterative depth-first traversal



An iterative implementation can use a stack

```
void Graph::depth_first_traversal( Vertex *first ) const {
    unordered_map<Vertex *, int> hash;
    hash.insert( first );
    std::stack<Vertex *> stack;
    stack.push( first );


    while ( !stack.empty() ) {
        Vertex *v = stack.top();
        stack.pop();
        // Perform an operation on v

        for ( Vertex *w : v->adjacent_vertices() ) {
            if ( !hash.member( w ) ) {
                hash.insert( w );
                stack.push( w );
            }
        }
    }
}
```

Example: BFS



Queue




Example

Performing a breadth-first traversal:

- Pop B and push D


Done: A, B



C	E	D			
---	---	---	--	--	--

Pop C and push F

A, B, C




E	D	F			
---	---	---	--	--	--

Example


Performing a breadth-first traversal:

- Pop E and push G,
- Done: A, B, C, E



Pop D


A, B, C, E, D



Example


Performing a breadth-first traversal:

- Pop F
- Done: A, B, C, E, D, F



G	H				
---	---	--	--	--	--

- Pop G, push I
A, B, C, E, D, F, G





H	I				
---	---	--	--	--	--

Example



Performing a breadth-first traversal:

- Pop H
- Done: A, B, C, E, D, F, G, H



Pop I

A, B, C, E, D, F, G, H, I








Figure 22.3 The operation of BFS on an undirected graph. Tree edges are shown shaded as they are produced by BFS. Within each vertex u is shown $d[u]$. The queue Q is shown at the beginning of each iteration of the while loop of lines 10–18. Vertex distances are shown next to vertices in the queue.

Example

Performing a recursive depth-first traversal:

- Insert A: Visited: A, B
Stack: A, B

Examine B: Visited A, B, C,
Stack A, B, C





Example

Performing a recursive depth-first traversal:

- Examine C: Visited: A, B, C, D
Stack: A, B, C, D

Pop D: Visited A, B, C, D
Stack A, B, C





Example

Performing a recursive depth-first traversal:

- Examine C: Visited: A,B,C,D,E
Stack: A,B,C,E

Examine E: Visited A,B,C,D,E,G
Stack A,B,C,E,G





Example

Performing a recursive depth-first traversal:

- Examine G: Visited: A,B,C,D,E,G,I
Stack: A,B,C,E,G,I

Examine I: Visited A,B,C,D,E,G,I,H
Stack A,B,C,E,G,I,H



Example

Performing a recursive depth-first traversal:

– Pop H: Visited: A,B,C,D,E,G,I,H

Stack: A,B,C,E,G,I

Pop G, then E



Inspect C: Visited A,B,C,D,E,G,I,H,F

Stack: A,B,C,F

Pop I: Visited A,B,C,D,E,G,I,H

Stack A,B,C,E,G,

Pop F, C, B, A. Done.




Comparison


The order in which vertices can differ greatly

- An iterative depth-first traversal may also be different again

A, B, C, E, D, F, G, H, I



A, B, C, D, E, G, I, H, F






Figure 22.4 The progress of the depth-first-search algorithm DFS on a directed graph. As edges are explored by the algorithm, they are shown as either shaded (if they are tree edges) or dashed (otherwise). Nontree edges are labeled B, C, or F according to whether they are back, cross, or forward edges. Vertices are timestamped by discovery time/finishing time.

Applications

Applications of tree traversals include:

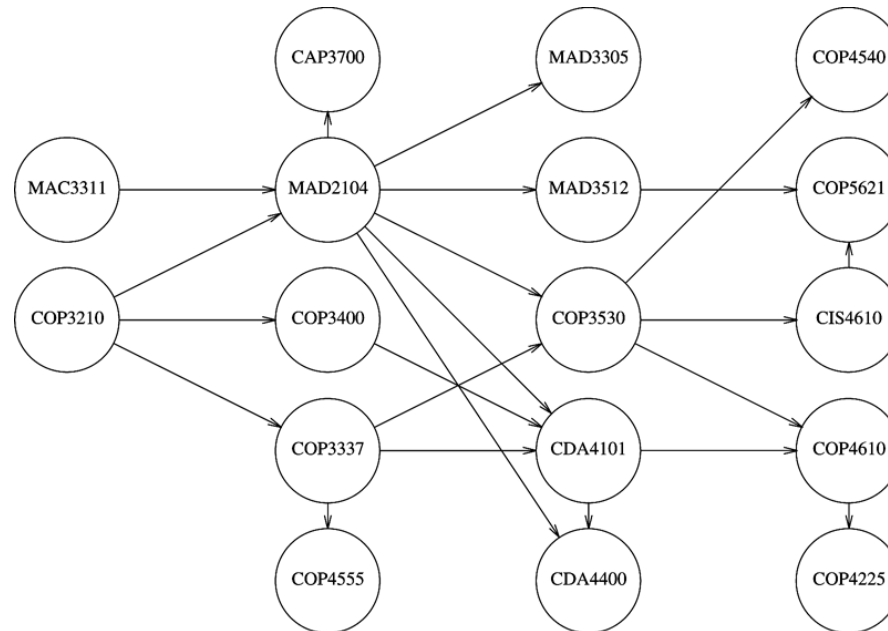
- Determining connectedness and finding connected sub-graphs
- Determining the path length from one vertex to all others
- Testing if a graph is bipartite
- Branch and bound search
- Topological Sort
- ...

Topological Sort

- Application
 - Given a number of tasks, there are often a number of constraints between the tasks:
 - task A must be completed before task B can start
 - These tasks together with the constraints form a directed acyclic graph
 - A topological sort of the graph gives an order in which the tasks can be scheduled while still satisfying the constraints


Topological Sort

- Course prerequisite structure represented in an acyclic graph



Topological Sort


- Topological Sort (Relabel): Given a directed acyclic graph (DAG), relabel the vertices such that every directed edge points from a lower-numbered vertex to a higher numbered one



Algorithm Using Priority

Determine the indegree of each vertex

List is set of vertices with indegree 0




next

0

Vertex
Indegree

1	2	3	4	5	6	7	8
2	2	3	2	1	1	0	2



Heap


155

Select a node from LIST

Select a vertex from LIST and delete it.

For all vertices in outlist of vertex, reduce in degree by 1

update LIST




next

1

Node
Indegree


Node	1	2	3	4	5	6	8
Indegree	2	2	3	1	0	1	2



LIST


156

Delete 5




next 2

Node	1	2	3	4		6	8
Indegree	2	1	3	0		0	2



157
LIST


Delete 6




next 3

Node	1	2	3	4
Indegree	1	0	3	0

8
2




Delete 2




next 4

Node	1	3	4
Indegree	0	3	0

8
2



Delete 1




next

5

Node
Indegree

3	4
2	0

8
1




LIST

160

Delete 4

next := next +1
order(i) := next;
update
indegrees
update LIST




next 6

Node
Indegree

3
1


8
0



LIST

Delete 8


next := next + 1
order(i) := next;
update
indegrees
update LIST



next 7

Node
Indegree


3
0



LIST

162

Delete 3




Node
Indegree

List is empty.

The algorithm
terminates with a
topological order of the
nodes

next

8




LIST

163

Complexity of Topological Sort

- Computing indegrees: $O(\#E)$
 - Recursively updating indegrees: One operation per edge, $O(\#E)$
 - So, algorithm is $O(\#E)$, using an array of $O(\#V)$
- Alternative algorithm: **recursive** DFS ($O(\#E)$)
 - Order in which vertices are completed is reverse order of a topological sort!



Stack: 1 2 4 Completed: 4

Stack: 1 2 3 Completed: 3

Stack: 1 2 5 Completed: 5


Stack: 1 2 Completed: 2

Stack: 1 Completed: 1

Stack: 6 Completed: 6

Reverse Order: 6,1,2,5,3,4

DAG Topological Sort (Execution order)



Biconnectivity


- Connected **undirected** graph is **biconnected** if there are no nodes whose removal disconnects the graph
- Nodes whose removal disconnect the graph are known as **articulation points**
- DFS can be used to find articulation points:
 - Algorithm: Number nodes in Depth First Search order, in the order in which they are inserted into the execution stack of the recursive Depth-First Search. This creates a spanning tree in the graph. Call this number $NUM(n)$ for node n

Biconnectivity

- Algorithm:
 - Number vertices in DFS order, in the order in which they are marked as visited in DFS. Call this number $NUM(n)$ for vertex n
 - For each vertex n , compute the lowest numbered vertex which is reachable from node n by following down the tree 0 or more steps and using a single edge which is not on the tree. Call that number $LOW(n)$
 - Find articulation points as follows:
 - a root of a DFS tree is an articulation point if and only if it has more than one child.
 - Any vertex n is an articulation point if and only if there is a child m of n in the DFS tree such that $LOW(m) \geq NUM(n)$

Biconnectivity

- For each vertex n , compute the lowest numbered vertex which is reachable from node n by following down the tree 0 or more steps and using a single edge which is not on the tree. Call that number $LOW(n)$




Low(A) = 1
Low(B) = 7
Low(C) = 6
Low(F) = 1 (path F-E-G-A)
Low(E) = 1
Low(D) = 2 (can't use 2 back edges)
Low(G) = 1

Biconnectivity

- Why does this work?
 - **Property 1:** The lowest numbered vertex reachable from n with one backward arc must be an ancestor of n in DFS tree
 - DFS algorithm guarantees that edges not included in the spanning tree cannot cross between branches.
 - **Property 2:** If, for all children n_1 of n in DFS tree, $LOW(n_1) < NUM(n)$, then there is path from each n_1 to an ancestor of n which bypasses n (because it is not in the DFS tree) \implies removing n leaves its ancestors and successors connected, so n cannot be an articulation point
 - **Property 3:** If there is a child n_1 of n for which $LOW(n_1) \geq NUM(n)$, then there is no path from n_1 to any ancestor of node n which does not go through $n \implies$ removing n disconnects $n_1 \implies n$ is articulation point

Example



NUM(1) = 1

NUM(2) = 2

NUM(3) = 9

NUM(4) = 8

NUM(5) = 3

NUM(6) = 7

NUM(7) = 6

NUM(8) = 5

NUM(9) = 4

LOW(1) = 1

LOW(2) = 1

LOW(3) = 1

LOW(4) = 1

LOW(5) = 1


LOW(6) = 1

LOW(7) = 1

LOW(8) = 1

LOW(9) = 1

DFS tree




NO ARTICULATION POINTS

Strong Connectivity

- A graph is strongly connected if every vertex can be reached from every other vertex
 - Would like to detect if a graph is strongly connected
- Lemma. Let s be any node. G is strongly connected iff every node is reachable from s , and s is reachable from every node.
- Algorithm: Pick any s .
 - Run BFS from s and verify all vertices can be reached
 - Reverse all edges in E , run BFS and verify all vertices can be reached
 - Complexity $O(\#E + \#V)$ (2 BFS).

Strongly-Connected Components



- A graph is strongly connected if every vertex can be reached from every other vertex
 - A strongly-connected component of a graph is a subgraph that is strongly connected
 - Would like to identify strongly-connected components of a graph
 - Can be used to identify weaknesses in a network
 - General approach: Perform two DFSs




Strongly-Connected Components

- Kosaraju's Algorithm
 - Perform DFS on graph $G = (V, E)$,
 - Number vertices according to their finishing time in DFS of G
 - Perform DFS on $G_r = (V, E_r)$, where E_r are reverse of edges in E , selecting nodes in decreasing order of finishing time in previous DFS
 - Strongly connected components = reachable trees obtained in last DFS


Example



Example



Reverse graph with distance labels



Reverse graph reachable trees

Strongly-Connected Components

- o Correctness
 - o If v and w are in a strongly-connected component
 - o Then there is a path from v to w and a path from w to v
 - o Therefore, there will also be a path between v and w in G and G^r
- o Running time
 - o Two executions of DFS
 - o $O(|E|+|V|)$