

EC504 ALGORITHMS AND DATA STRUCTURES
FALL 2020 MONDAY & WEDNESDAY
2:30 PM - 4:15 PM

Prof: David Castañón, dac@bu.edu

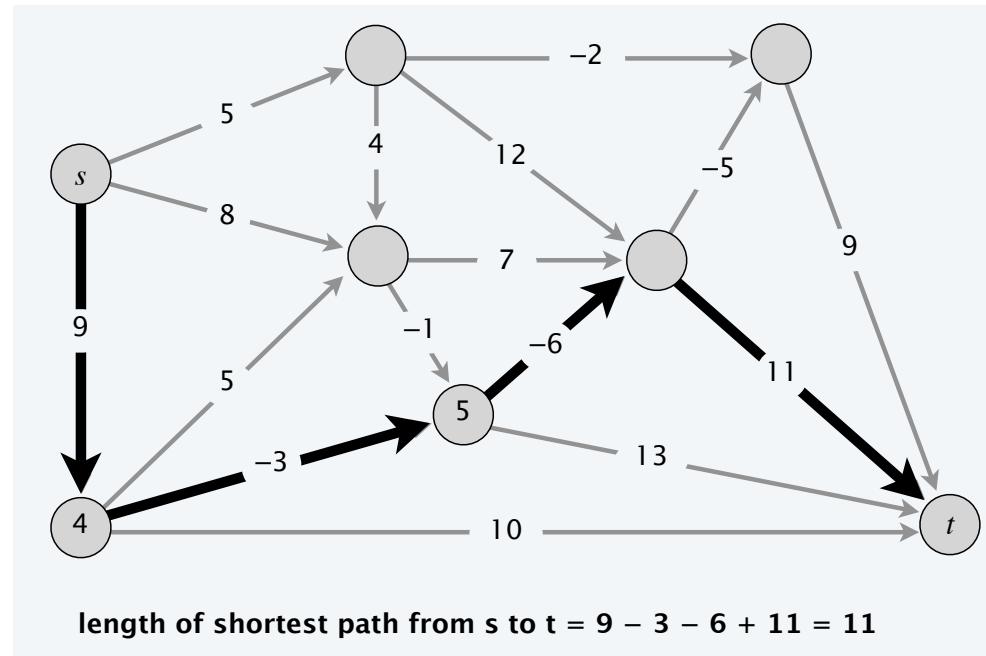
GTF: Mert Toslali, toslali@bu.edu

Haoyang Wang: haoxyangw@bu.edu

Christopher Liao: cliao25@bu.edu

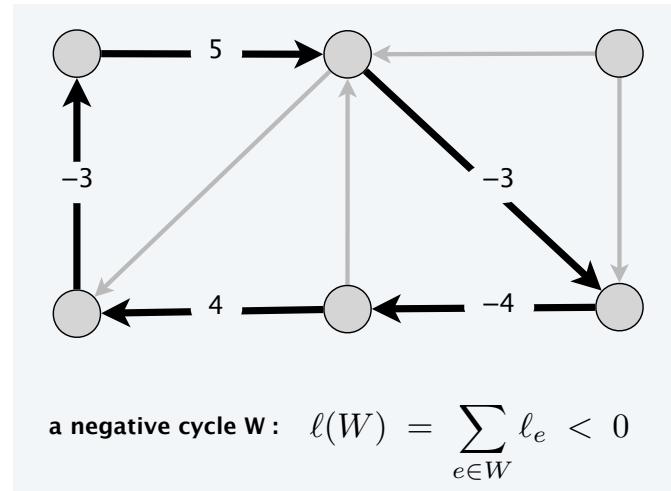
Shortest paths: Graphs with Negative Weights

- ▶ Need different algorithm
- ▶ Dijkstra's algorithm: "label-setting"
 - ▶ Once vertex exits priority queue, its distance label never changes
- ▶ Different class of algorithms: "label-correcting"
 - ▶ Exploit Bellman's principle of optimality
 - ▶ Will handle some cases with negative weights



Limitations: Negative Cycles

- A negative cycle is a directed cycle where the sum of the weights on the edges is negative
- **Lemma:** if some path from s to t contains a negative cycle, there does not exist a shortest path from s to t
 - Can go around the negative cycle an infinite amount of time...
- **Lemma:** if there exist no negative cycles in directed graph, then there exists a shortest path from s to t that is simple (no cycles) and has $< |V|$ edges
 - Consider minimum number of edges among all shortest paths from s to t
 - If numbers edges $\geq |V|$, then it must have a cycle
 - If this path has a cycle, can remove cycle and get one with less edges and same total weight



Dynamic Programming

- ▶ **Dynamic programming**: an algorithmic technique for solving an optimization problem involving sequences of decisions by
 - ▶ Breaking it down into simpler subproblems
 - ▶ Exploiting that the optimal solution to the overall problem depends upon the optimal solution to its subproblems
 - ▶ Bellman's principle of optimality: shortest path from s to t is composed of shortest paths to and from any intermediate vertices
- ▶ Idea: let's solve subproblem of finding shortest path with a maximum number of edges
 - ▶ Easy if we limit it to 1; if we extend it to $|V|-1$, we solved the full problem

Bellman-Ford Algorithm

- **Notation:** $D(k,v)$ is shortest distance from vertex s to vertex v on paths with less than or equal to k edges

Initial value: $D(0,v) = \infty$, $v \neq s$; $0, v = s$

Bellman equation $D(k, v) = \min \left\{ D(k - 1, v), \min_{(u,v) \in E} [D(k - 1, u) + w(u, v)] \right\}$ (DP)

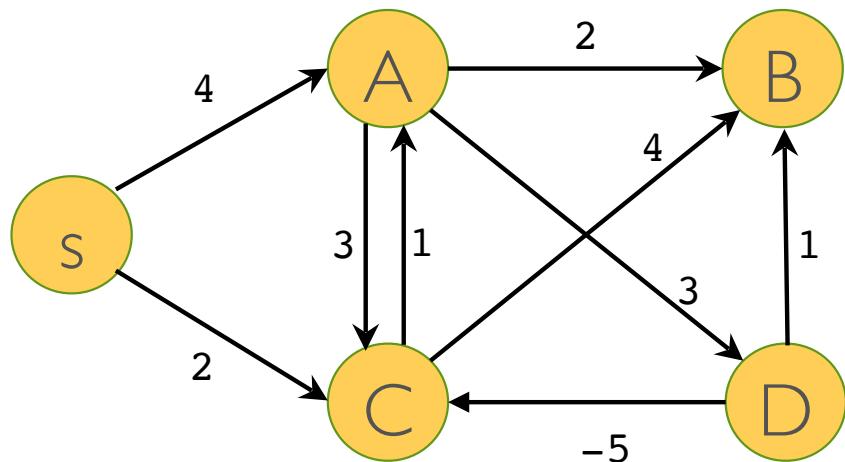
- Bellman-Ford Algorithm
 - For each $v \in V$, set $D(0,v) = \infty$, $v \neq s$; $0, v = s$; set $\text{pred}[v] = \text{null}$
 - For $k = 1$ to $|V| - 1$:
 - For each $v \in V$, set $D(k, v) = D(k - 1, v)$
 - For each $(u, v) \in E$:
 - if $D(k, v) > D(k - 1, u) + w(u, v)$:
 - Set $D(k, v) = D(k - 1, u) + w(u, v)$, set $\text{pred}[v] = u$

Bellman-Ford Algorithm Properties

- $D(k, v)$ is the shortest distance for a path of k or less edges from s to v
 - Prove by induction: obviously true at $k = 0$. Assume true for $k-1$. Assume $D(k, v) = D(k - 1, v)$. Then, from the iteration, no path with k vertices is shorter than $D(k, v)$. Note that $D(k - 1, u) + w(u, v)$ is the shortest distance of all paths of length k or less that end in v and have next to last vertex u . This proves the result, because $D(k, v) = \min_{(u,v) \in E} D(k - 1, u) + w(u, v)$
- At each iteration, $D(k, v)$ is upper bound on the shortest path distance from s to v
- $D(k, v)$ is non-increasing in k
- If $D(k - 2, v) = D(k - 1, v)$, one does not have to scan edges from node v for $D(k, v)$
 - Best path to v is less than $k-1$ edges
- Unlike Dijkstra's algorithm, each edge is examined more than once; once per k
- Brute force complexity: $|V|$ iterations, each of which is $O(|E|) \rightarrow$ complexity $O(|V||E|)$

Bellman-Ford Algorithm Properties

- Algorithm converges to shortest paths provided no negative cycles exist in graph
- If negative cycles exist, predecessor graph will have cycles!
- Distances continue decreasing after $|V|-1$ iterations (neg. cycle detection)



	D(s)	D(A)	D(B)	D(C)	D(D)
k=1	0	4	Inf	2	Inf
k=2	0	3	6	2	7
k=3	0	3	5	2	6
k=4	0	3	5	1	6

Final pred array: $\text{pred}[s] = \text{null}$, $\text{pred}[A] = C$, $\text{pred}[B] = A$, $\text{pred}[C] = D$, $\text{pred}[D] = A$

Better Bellman-Ford Implementation

- Use a queue of vertices where distances have changed
 - Avoid looking at edges out of vertices where distances have not changed
- Algorithm
 - Initialize $D[s] = 0; D[v] = \infty, v \neq s;$ $\text{pred}[v] = \text{null}$
 - Insert s into queue $Q;$ mark $\text{inqueue}[s] = \text{true},$ mark $\text{inqueue}[v] = \text{false}, v \neq s$
 - While Q is not empty:
 - Select u out of queue, mark $\text{inqueue}[u] = \text{false}$
 - For each edge (u,v) in $E:$
 - If $D[v] > D[u] + w(u,v):$
 - Set $D[v] = D[u] + w(u,v),$ set $\text{pred}[v] = u$
 - if $\text{inqueue}[v] = \text{false},$ add v to $Q,$ mark $\text{inqueue}[v] = \text{true}$

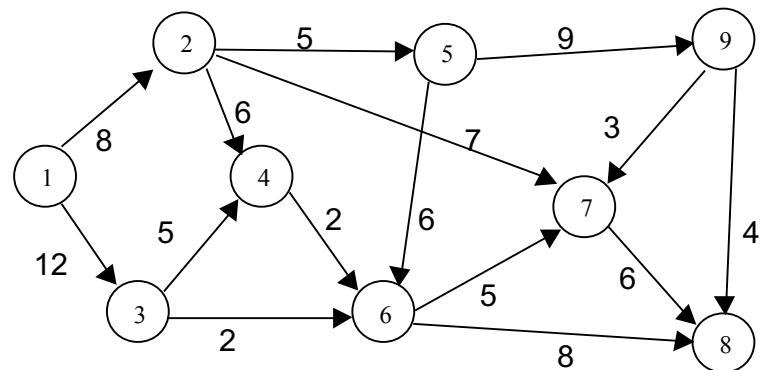
Better Bellman-Ford Implementation

- Many variations possible
 - Have not specified in what order to insert vertices into queue Q
 - Have not specified in what order to extract vertices from Q
 - Algorithm works anyway, as long as there are no negative weight cycles
- Detecting negative cycles
 - A negative cycle exists if and only if a vertex enters the Q $|V|$ times
- Simplest implementation: use queue in FIFO order
- More interesting implementations: insert vertices with shorter distances $D[v]$ in the front of the queue
- Convergence happens because distance labels $D[v]$ decrease monotonically, and are bounded below
- Bellman-Ford and variations are known as label-correcting algorithms: distance labels $D[v]$ are not set when v leaves the Queue
 - Less overhead than Dijkstra; can be empirically faster

Example

- Source node 1

If $D[v] > D[u] + w(u, v)$, set $D[v] = D[u] + w(u, v)$

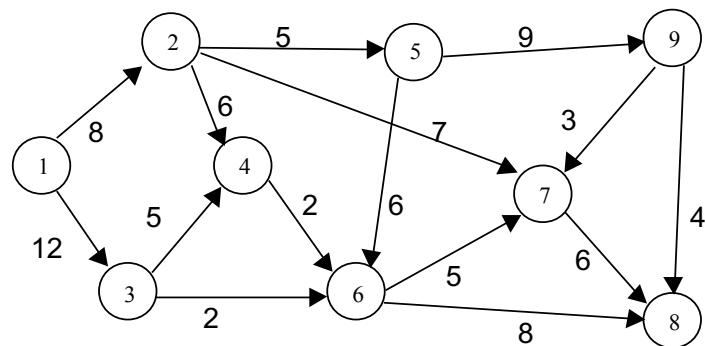


	D[]	pred[]	Q []
1	0	Null	1
2	Inf	Null	
3	Inf	Null	
4	Inf	Null	
5	Inf	Null	
6	Inf	Null	
7	Inf	Null	
8	Inf	Null	
9	Inf	Null	

Pop 1
➡

	D[]	pred[]	Q []
1	0	Null	
2	8	1	1
3	12	1	2
4	Inf	Null	
5	Inf	Null	
6	Inf	Null	
7	Inf	Null	
8	Inf	Null	
9	Inf	Null	

Example - 2



If $D[v] > D[u] + w(u, v)$, set $D[v] = D[u] + w(u, v)$

	D[]	pred[]	Q []
1	0	Null	
2	8	1	1
3	12	1	2
4	Inf	Null	
5	Inf	Null	
6	Inf	Null	
7	Inf	Null	
8	Inf	Null	
9	Inf	Null	

Pop 2



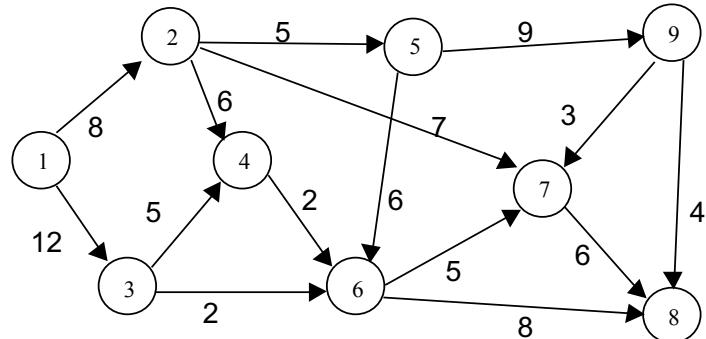
	D[]	pred[]	Q []
1	0	Null	
2	8	1	
3	12	1	1
4	14	2	2
5	13	2	3
6	Inf	Null	
7	15	2	4
8	Inf	Null	
9	Inf	Null	

Pop 3



	D[]	pred[]	Q []
1	0	Null	
2	8	1	
3	12	1	
4	14	2	1
5	13	2	2
6	14	3	4
7	15	2	3
8	Inf	Null	
9	Inf	Null	

Example - 3



If $D[v] > D[u] + w(u, v)$, set $D[v] = D[u] + w(u, v)$

	$D[]$	$\text{pred}[]$	$Q[]$
1	0	Null	
2	8	1	
3	12	1	
4	14	2	1
5	13	2	2
6	14	3	4
7	15	2	3
8	Inf	Null	
9	Inf	Null	

Pop 4



	$D[]$	$\text{pred}[]$	$Q[]$
1	0	Null	
2	8	1	
3	12	1	
4	14	2	
5	13	2	1
6	14	3	3
7	15	2	2
8	Inf	Null	
9	Inf	Null	

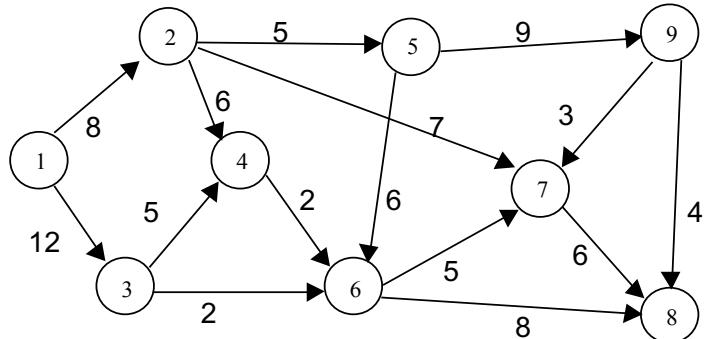
i-to

	$D[]$	$\text{pred}[]$	$Q[]$
1	0	Null	
2	8	1	
3	12	1	
4	14	2	
5	13	2	
6	14	3	2
7	15	2	1
8	Inf	Null	
9	22	5	3

Pop 5



Example - 4



If $D[v] > D[u] + w(u, v)$, set $D[v] = D[u] + w(u, v)$

	D[]	pred[]	Q []
1	0	Null	
2	8	1	
3	12	1	
4	14	2	
5	13	2	
6	14	3	2
7	15	2	1
8	Inf	Null	
9	22	5	3

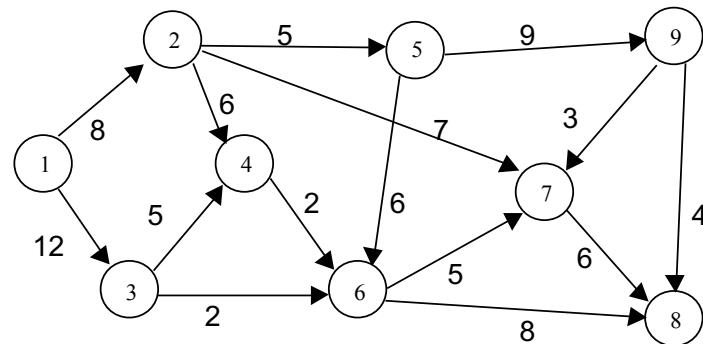
Pop 7
→

	D[]	pred[]	Q []
1	0	Null	
2	8	1	
3	12	1	
4	14	2	
5	13	2	
6	14	3	1
7	15	2	
8	21	7	3
9	22	5	2

Pop 6
→

	D[]	pred[]	Q []
1	0	Null	
2	8	1	
3	12	1	
4	14	2	
5	13	2	
6	14	3	
7	15	2	
8	21	7	2
9	22	5	1

Example - 5



If $D[v] > D[u] + w(u, v)$, set $D[v] = D[u] + w(u, v)$

	D[]	pred[]	Q []
1	0	Null	
2	8	1	
3	12	1	
4	14	2	
5	13	2	
6	14	3	
7	15	2	
8	21	7	2
9	22	5	1

Pop 9
→

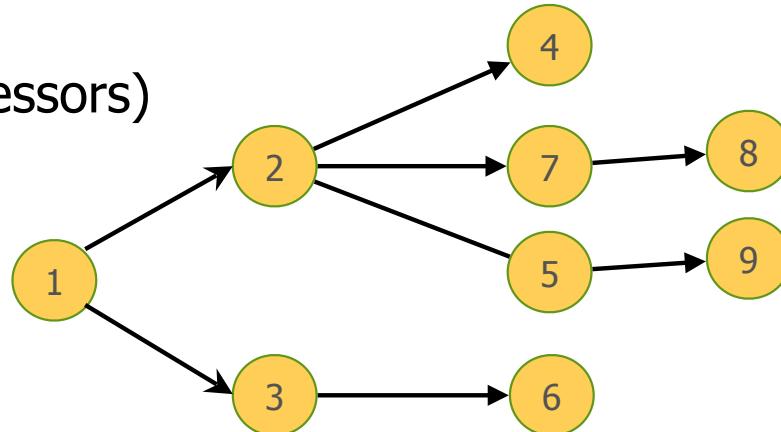
	D[]	pred[]	Q []
1	0	Null	
2	8	1	
3	12	1	
4	14	2	
5	13	2	
6	14	3	
7	15	2	
8	21	7	1
9	22	5	

→

	D[]	pred[]	Q []
1	0	Null	
2	8	1	
3	12	1	
4	14	2	
5	13	2	
6	14	3	
7	15	2	
8	21	7	
9	22	5	

Example - 6

- ▶ Observations
 - ▶ Each vertex left queue only once, so edges out of each vertex were scanned only once
 - ▶ With simple data structure, this is much faster than Dijkstra
 - ▶ Compared with full Bellman-Ford, each edge would be scanned 8 times ($|V| - 1$)
- ▶ Final shortest path tree (use predecessors)



Bellman-Ford vs Dijkstra

- Dijkstra's algorithm
 - At each step, finds the next vertex that is closest to the start vertex
 - Grows the shortest path tree in order of shortest distance
 - Greedy: no backtracking
 - $O(|E| + |V| \log(|V|))$ with complex data structures
 - Requires non-negative weights
- Bellman-Ford
 - At each step, finds shortest distance to vertices over paths with less than k edges
 - Grows the shortest path tree in order of minimum number of edges from start vertex
 - Backtracks, can scan the same vertex multiple times
 - $O(|E||V|)$ worst case with very simple data structures
 - Can handle negative weights, but no negative cycles
 - Will work asynchronously, perfect for parallel applications

New Variation: All Pairs Shortest Paths

Input: Directed graph $G = (V, E)$, where $V = \{1, 2, \dots, n\}$, with edge-weight function $w : E \rightarrow \mathbb{R}$

- Weights may be negative, but no negative cycles

Output: $n \times n$ matrix of shortest-path lengths $D(i, j)$ for all $i, j \in V$ (and routes)

Possible solution:

- Run Bellman-Ford once from each vertex as starting node
- Time $O(|V|^2 E)$
- Dense graph ($\Theta(n^2)$ edges) $\Rightarrow O(n^4)$ time in the worst case

New Variation: All Pairs Shortest Paths

Input: Directed graph $G = (V, E)$, where $V = \{1, 2, \dots, n\}$, with edge-weight function $w : E \rightarrow \mathbb{R}$

- Assume all weights are non-negative

Possible solution:

- Run Dijkstra once from each vertex as starting node
- Time $O(|V||E| + |V|^2 \log(|V|))$
- Dense graph ($\Theta(n^2)$) $\Rightarrow O(n^3)$ time in the worst case with array priority queue

Better Approach: Use Dynamic Programming

- ▶ Concept
 - ▶ Let $D(i,j)$ be the shortest distance from vertex i to vertex j .
 - ▶ Assume $D(j,j) = 0$
 - ▶ Then, shortest path from i to j must have at least one vertex before j (could be i)
- ▶ Bellman's principle of optimality: shortest path is composed of shortest paths

$$D(i,j) = \min_{k: (k,j) \in E} [D[i,k] + w(k,j)]$$

$$D(i,j) = \min_{k: k \in V} [D[i,k] + D[k,j]]$$

Floyd-Warshall Algorithm

- Initialize $n \times n$ weighted adjacency matrix A and initial distance matrix $D^{(0)}$ as

$$A(i, j) = \begin{cases} 0 & \text{if } i = j \\ w(i, j) & \text{if } (i, j) \in E, \\ \infty & \text{if } (i, j) \notin E \end{cases}$$
$$D^{(0)}(i, j) = \begin{cases} 0 & \text{if } i = j \\ \infty & \text{if } i \neq j \end{cases}$$

- Note that $D^{(0)}$ is the matrix of shortest distances among paths with zero edges!
- A simple algorithm

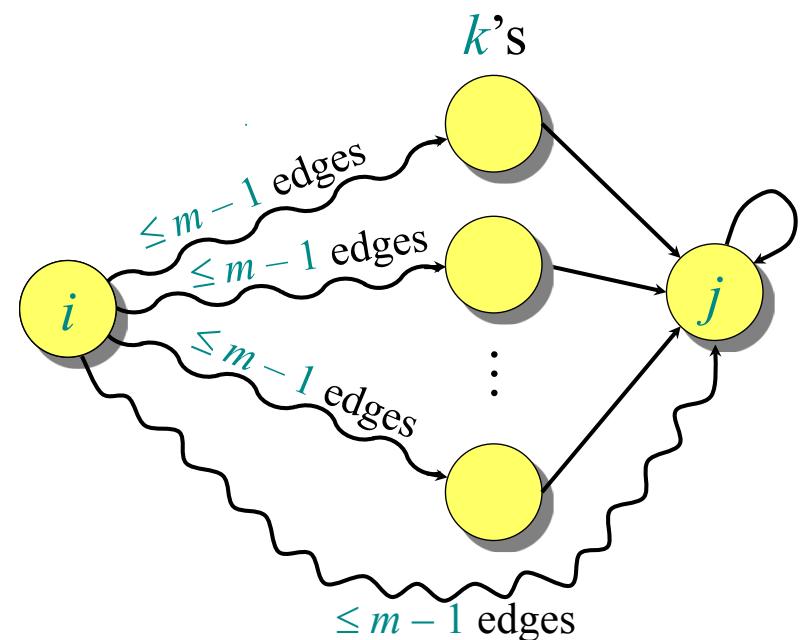
- For $m = 1, 2, \dots, n-1$:

$$D^{(m)}(i, j) = \min_{k \in V} [D^{(m-1)}(i, k) + A(k, j)]$$

- Note that $A(j, j) = 0$, so that $D^{(m)}(i, j) = D^{(m-1)}(i, j)$ is possible, and thus $D^{(m)}(i, j)$ is non-increasing

Floyd-Warshall Algorithm

- Claim: $D^{(m)}(i, j)$ is shortest distance from i to j among paths with m or less edges
 - Any path with less than or equal to m edges must have next to last vertex with shortest path less than $m-1$ edges
 - Principle of optimality



Floyd-Warshall Algorithm

- › Variation: faster algorithm with less memory

Initially $D(i, j) = \begin{cases} w(i, j) & \text{if } (i, j) \in E \\ \infty & \text{if } (i, j) \notin E \end{cases}$

- › For k in V :

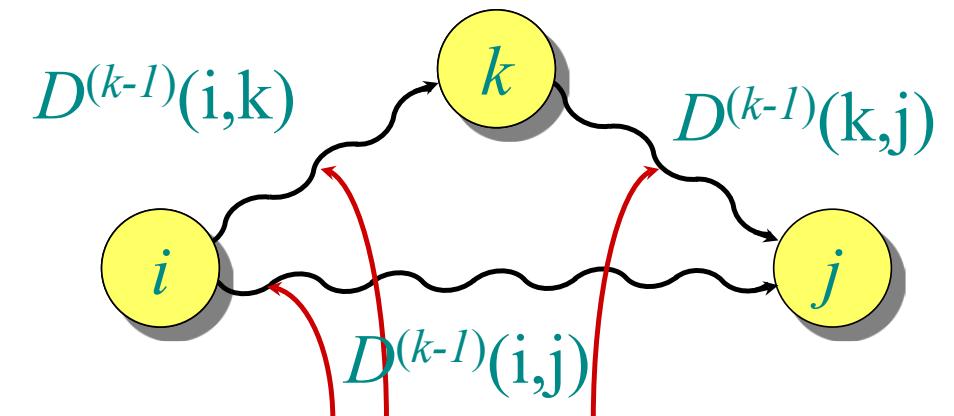
- › For i in V :

- › For j in V :

If $D(i, j) > D(i, k) + D(k, j)$:

set $D(i, j) = D(i, k) + D(k, j)$

- Complexity: $O(|V|^3)$



Floyd-Warshall Algorithm

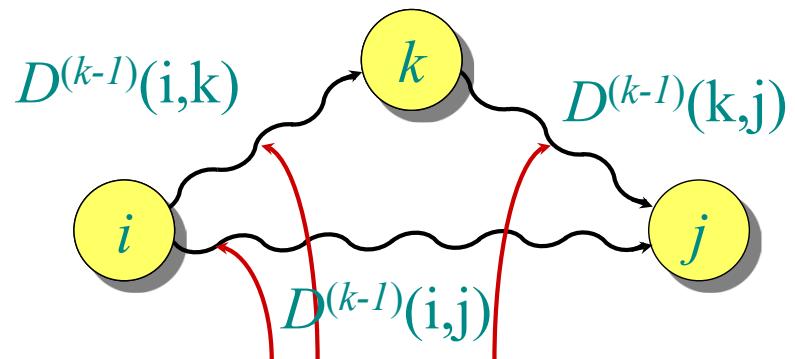
- Keeping track of paths

$$\text{Initially } pred(i, j) = \begin{cases} i & (i, j) \in E \\ -1 & (i, j) \notin E \end{cases} \quad D(i, j) = \begin{cases} w(i, j) & \text{if } (i, j) \in E \\ \infty & \text{if } (i, j) \notin E \end{cases}$$

- For k in V :
 - For i in V :
 - For j in V :

If $D(i, j) > D(i, k) + D(k, j)$:

set $D(i, j) = D(i, k) + D(k, j)$, $pred(i, j) = pred(k, j)$



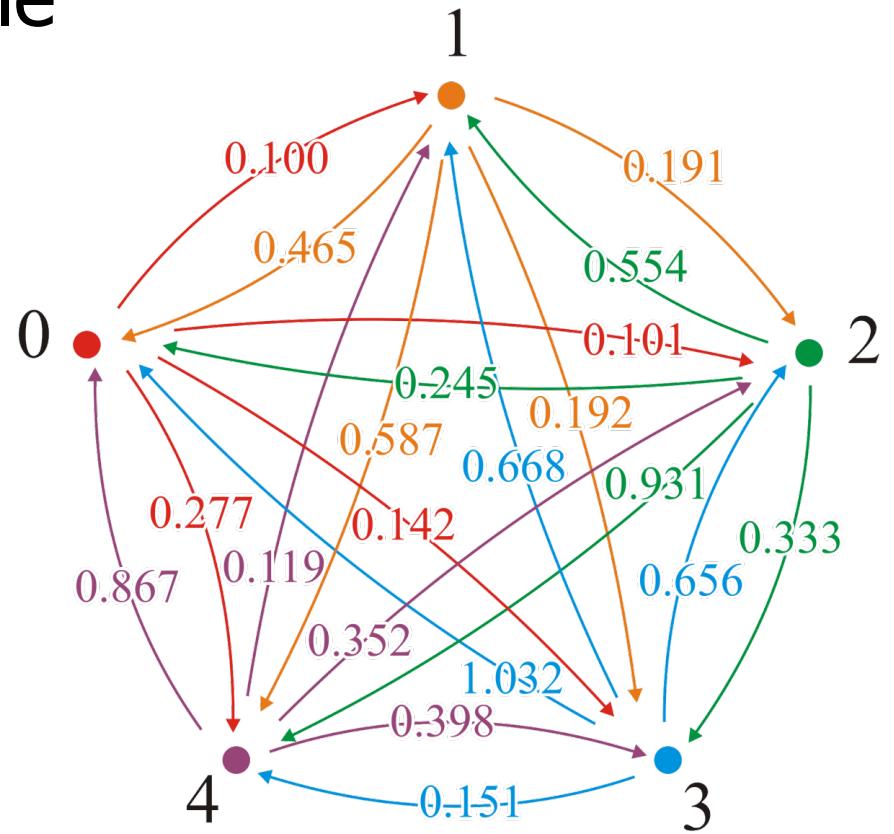
Example

- Weighted adjacency matrix

A =

0	0.1	0.101	0.142	0.277
0.465	0	0.191	0.192	0.587
0.245	0.554	0	0.333	0.931
1.032	0.668	0.656	0	0.151
0.867	0.119	0.352	0.398	0

Initially, $D = A$ (fully connected)



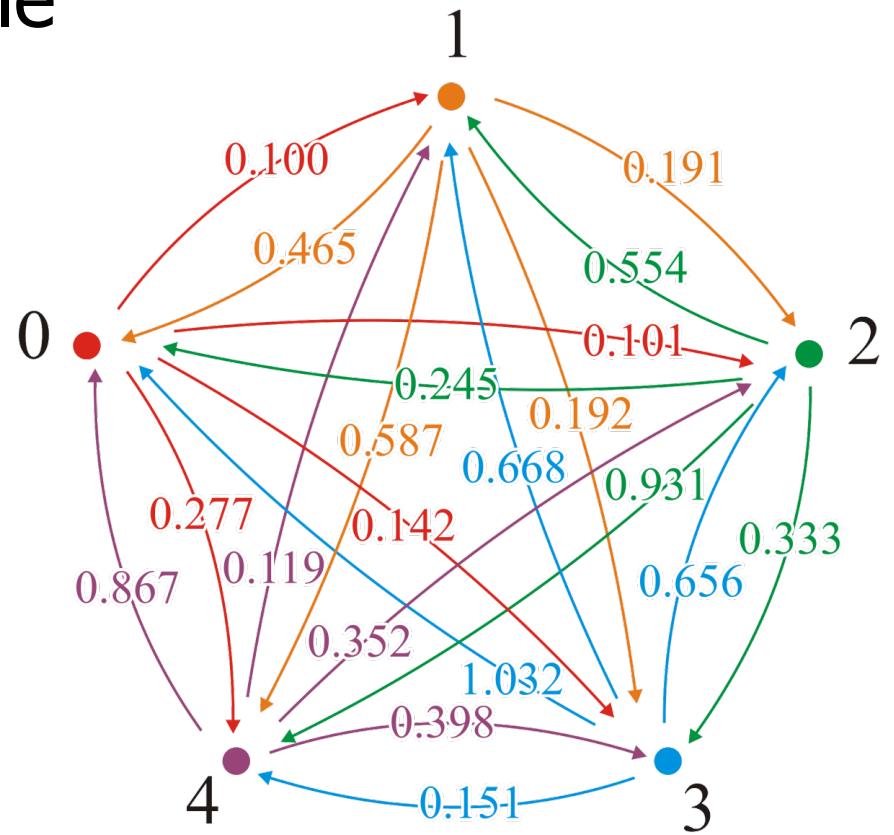
Example

- Iteration for $k = 0$

► D=

0	0.1	0.101	0.142	0.277
0.465	0	0.191	0.192	0.587
0.245	0.345	0	0.333	0.522
1.032	0.668	0.656	0	0.151
0.867	0.119	0.352	0.398	0

-1	0	0	0	0
1	-1	1	1	1
2	0	-1	2	0
3	3	3	-1	3
4	4	4	4	-1



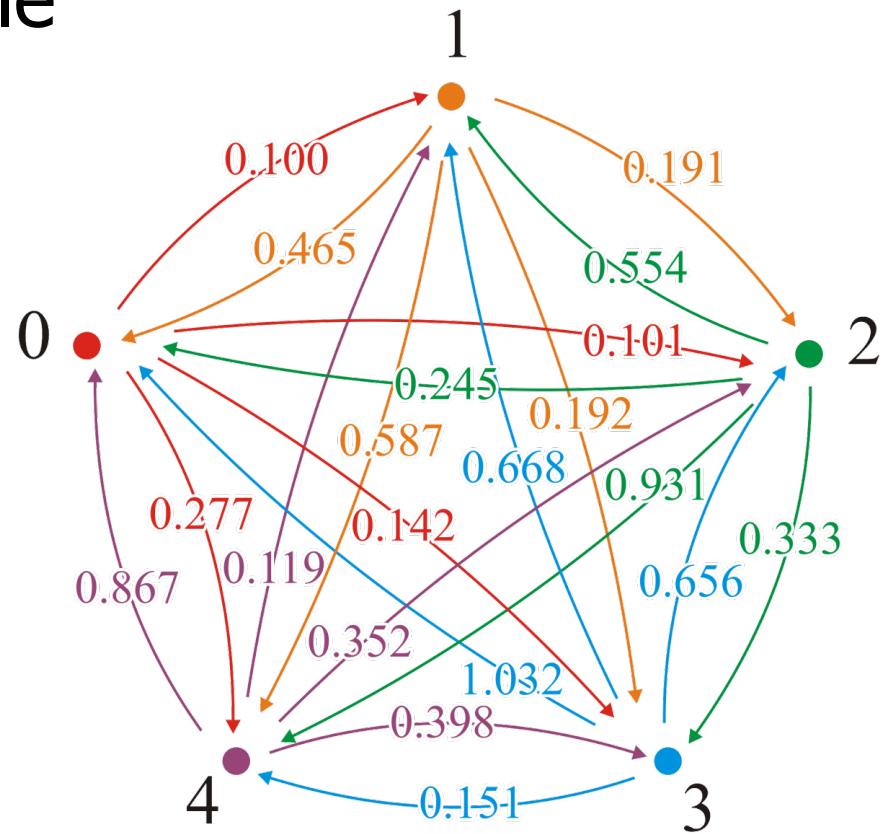
Example

- Iteration for $k = 1$

► D=

0	0.1	0.101	0.142	0.277
0.465	0	0.191	0.192	0.587
0.245	0.345	0	0.333	0.522
1.032	0.668	0.656	0	0.151
0.584	0.119	0.31	0.311	0

-1	0	0	0	0
1	-1	1	1	1
2	0	-1	2	0
3	3	3	-1	3
1	4	1	1	-1



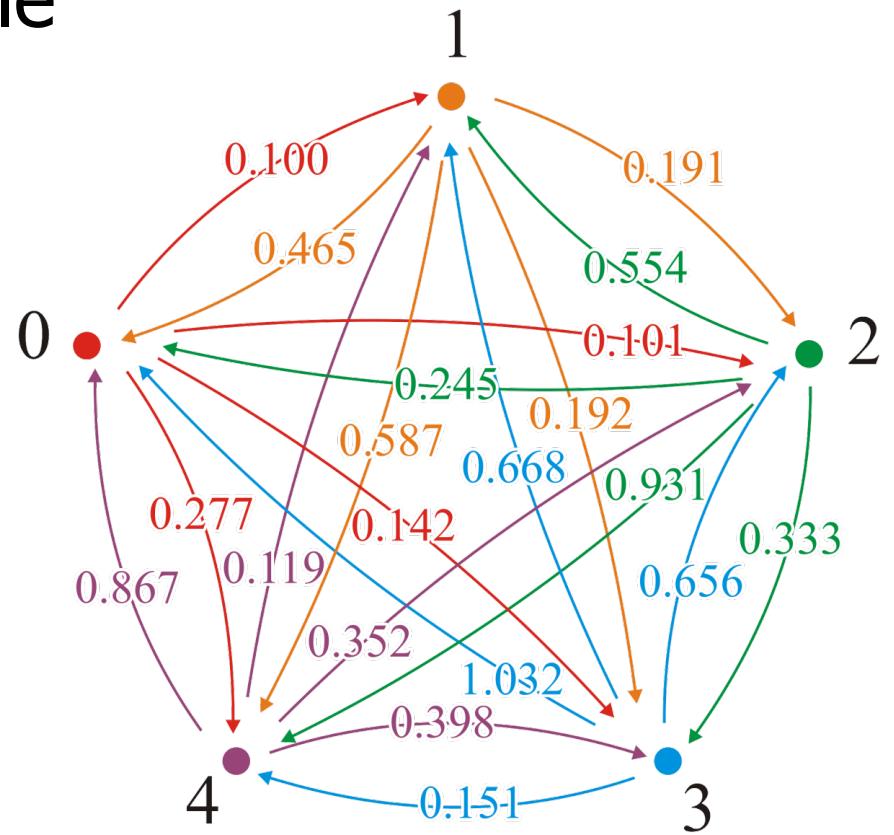
Example

- Iteration for $k = 2$

► D=

0	0.1	0.101	0.142	0.277
0.436	0	0.191	0.192	0.587
0.245	0.345	0	0.333	0.522
0.901	0.668	0.656	0	0.151
0.555	0.119	0.31	0.311	0

-1	0	0	0	0
2	-1	1	1	1
2	0	-1	2	0
2	3	3	-1	3
2	4	1	1	-1



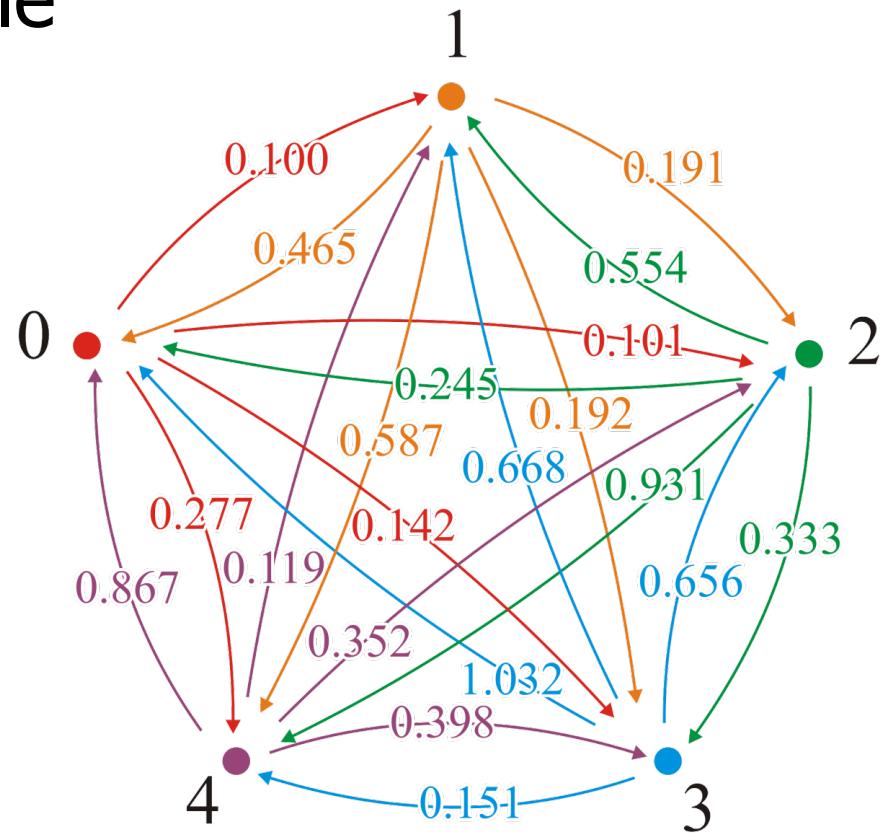
Example

- Iteration for $k = 3$

► D=

0	0.1	0.101	0.142	0.277
0.436	0	0.191	0.192	0.343
0.245	0.345	0	0.333	0.484
0.901	0.668	0.656	0	0.151
0.555	0.119	0.31	0.311	0

-1	0	0	0	0
2	-1	1	1	3
2	0	-1	2	3
2	3	3	-1	3
2	4	1	1	-1



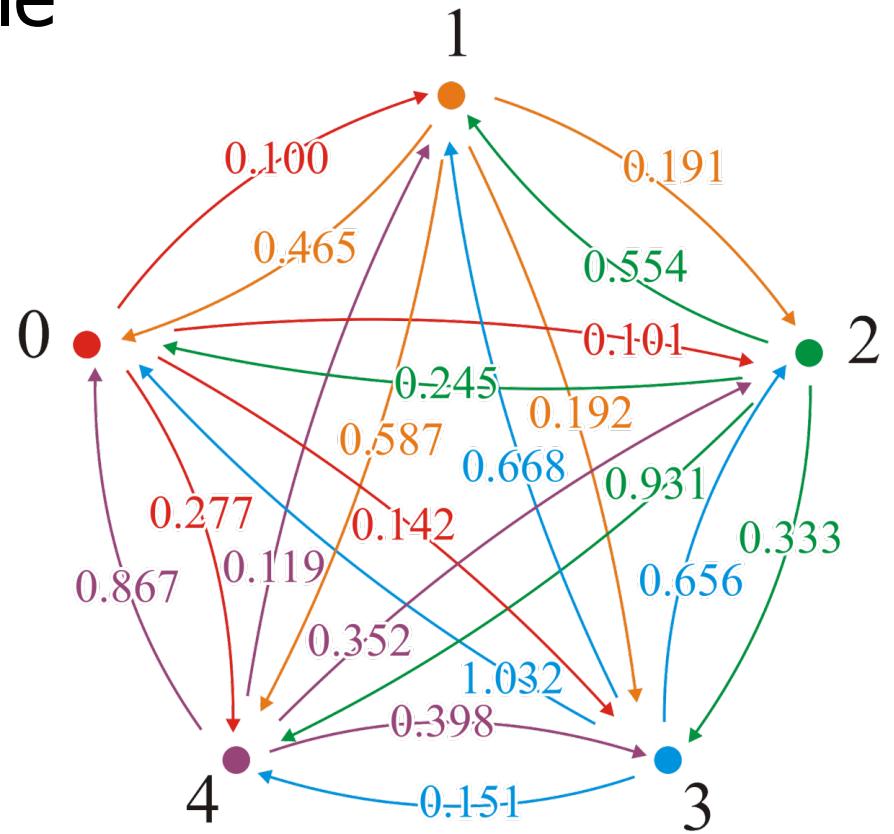
Example

- Iteration for $k = 4$

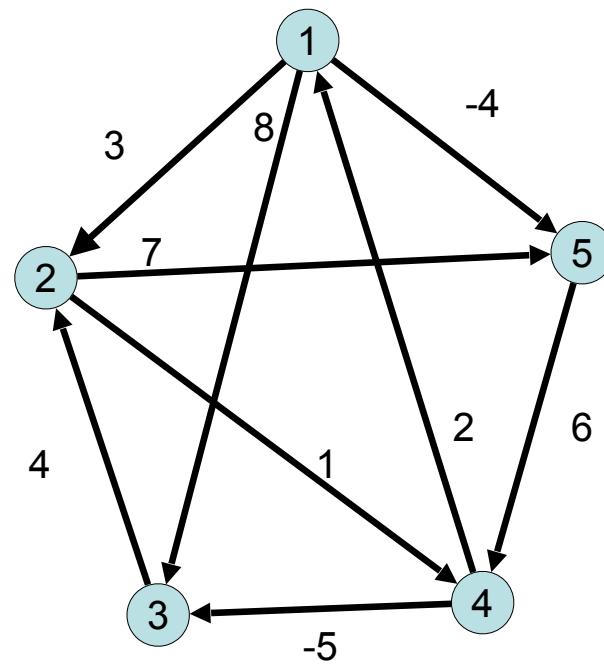
► D=

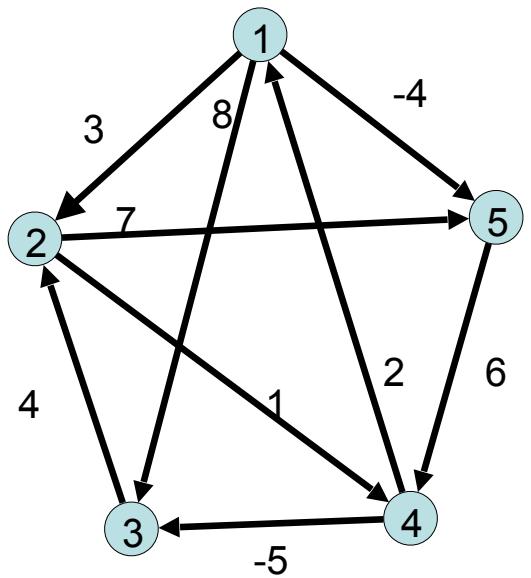
0	0.1	0.101	0.142	0.277
0.436	0	0.191	0.192	0.343
0.245	0.345	0	0.333	0.484
0.706	0.27	0.461	0	0.151
0.555	0.119	0.31	0.311	0

-1	0	0	0	0
2	-1	1	1	3
2	0	-1	2	3
2	4	1	-1	3
2	4	1	1	-1



Floyd-Warshall: Example 2

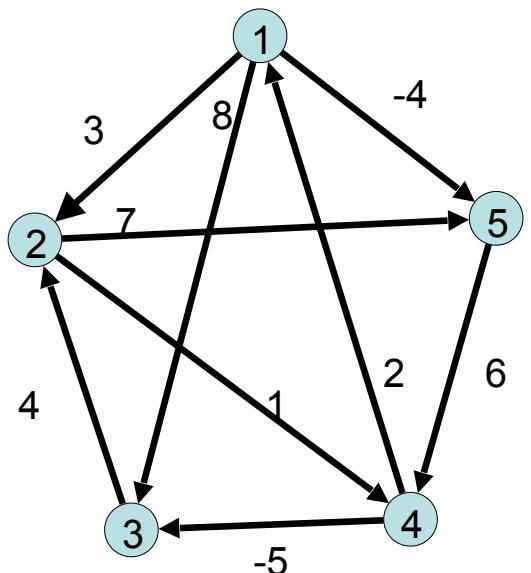




Floyd-Warshall

Initial distances and predecessors

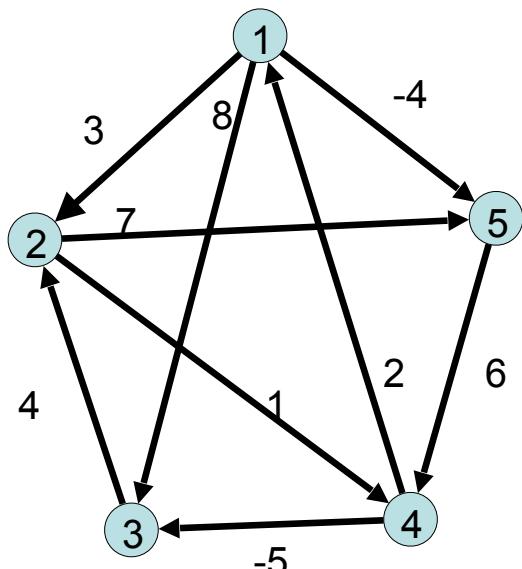
$$D^{(0)} = \begin{pmatrix} 0 & 3 & 8 & \infty & -4 \\ \infty & 0 & \infty & 1 & 7 \\ \infty & 4 & 0 & \infty & \infty \\ 2 & \infty & -5 & 0 & \infty \\ \infty & \infty & \infty & 6 & 0 \end{pmatrix} \quad \Pi^{(0)} = \begin{pmatrix} NIL & 1 & 1 & NIL & 1 \\ NIL & NIL & NIL & 2 & 2 \\ NIL & 3 & NIL & NIL & NIL \\ 4 & NIL & 4 & NIL & NIL \\ NIL & NIL & NIL & 5 & NIL \end{pmatrix}$$



Floyd-Warshall

After k = 1

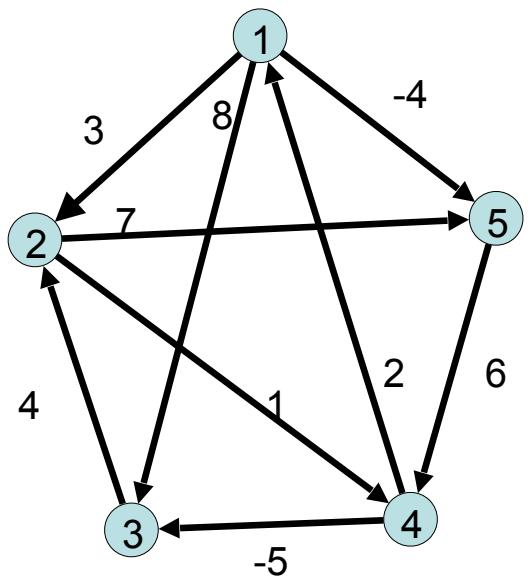
$$D^{(1)} = \begin{pmatrix} 0 & 3 & 8 & \infty & -4 \\ \infty & 0 & \infty & 1 & 7 \\ \infty & 4 & 0 & \infty & \infty \\ 2 & 5 & -5 & 0 & -2 \\ \infty & \infty & \infty & 6 & 0 \end{pmatrix} \quad \Pi^{(1)} = \begin{pmatrix} NIL & 1 & 1 & NIL & 1 \\ NIL & NIL & NIL & 2 & 2 \\ NIL & 3 & NIL & NIL & NIL \\ 4 & 1 & 4 & NIL & 1 \\ NIL & NIL & NIL & 5 & NIL \end{pmatrix}$$



Floyd-Warshall

After k = 2

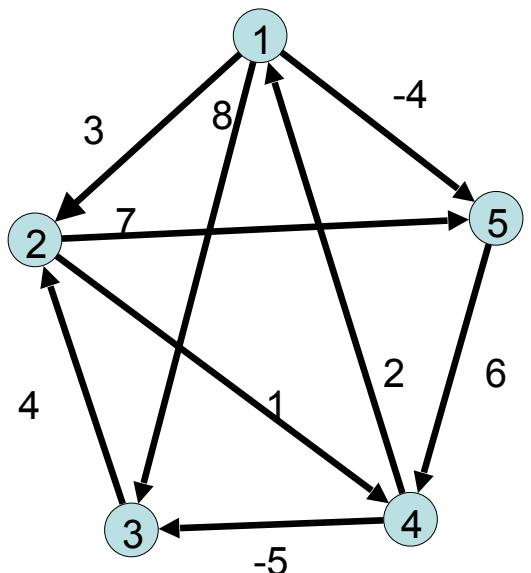
$$D^{(2)} = \begin{pmatrix} 0 & 3 & 8 & 4 & -4 \\ \infty & 0 & \infty & 1 & 7 \\ \infty & 4 & 0 & 5 & 11 \\ 2 & 5 & -5 & 0 & -2 \\ \infty & \infty & \infty & 6 & 0 \end{pmatrix} \quad \Pi^{(2)} = \begin{pmatrix} NIL & 1 & 1 & 2 & 1 \\ NIL & NIL & NIL & 2 & 2 \\ NIL & 3 & NIL & 2 & 2 \\ 4 & 1 & 4 & NIL & 1 \\ NIL & NIL & NIL & 5 & NIL \end{pmatrix}$$



Floyd-Warshall

After k = 3

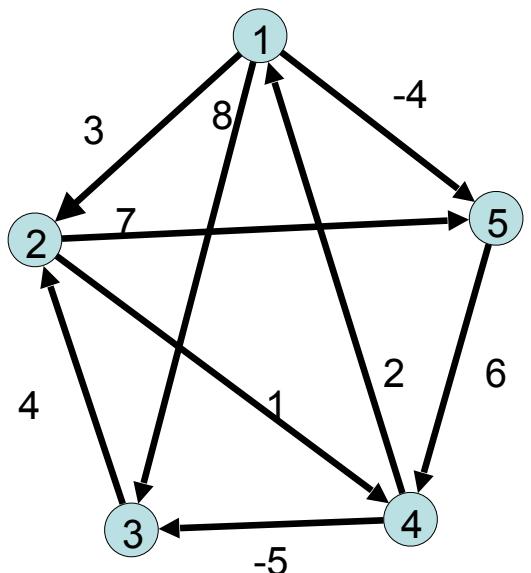
$$D^{(3)} = \begin{pmatrix} 0 & 3 & 8 & 4 & -4 \\ \infty & 0 & \infty & 1 & 7 \\ \infty & 4 & 0 & 5 & 11 \\ 2 & -1 & -5 & 0 & -2 \\ \infty & \infty & \infty & 6 & 0 \end{pmatrix} \quad \Pi^{(3)} = \begin{pmatrix} NIL & 1 & 1 & 2 & 1 \\ NIL & NIL & NIL & 2 & 2 \\ NIL & 3 & NIL & 2 & 2 \\ 4 & 3 & 4 & NIL & 1 \\ NIL & NIL & NIL & 5 & NIL \end{pmatrix}$$



Floyd-Warshall

After k = 4

$$D^{(4)} = \begin{pmatrix} 0 & 3 & -1 & 4 & -4 \\ 3 & 0 & -4 & 1 & -1 \\ 7 & 4 & 0 & 5 & 3 \\ 2 & -1 & -5 & 0 & -2 \\ 8 & 5 & 1 & 6 & 0 \end{pmatrix} \quad \Pi^{(4)} = \begin{pmatrix} NIL & 1 & 4 & 2 & 1 \\ 4 & NIL & 4 & 2 & 1 \\ 4 & 3 & NIL & 2 & 1 \\ 4 & 3 & 4 & NIL & 1 \\ 4 & 3 & 4 & 5 & NIL \end{pmatrix}$$



Floyd-Warshall

After k = 5 (final)

$$D^{(5)} = \begin{pmatrix} 0 & 1 & -3 & 2 & -4 \\ 3 & 0 & -4 & 1 & -1 \\ 7 & 4 & 0 & 5 & 3 \\ 2 & -1 & -5 & 0 & -2 \\ 8 & 5 & 1 & 6 & 0 \end{pmatrix} \quad \Pi^{(5)} = \begin{pmatrix} NIL & 3 & 4 & 5 & 1 \\ 4 & NIL & 4 & 2 & 1 \\ 4 & 3 & NIL & 2 & 1 \\ 4 & 3 & 4 & NIL & 1 \\ 4 & 3 & 4 & 5 & NIL \end{pmatrix}$$

Can We Do Better for Sparse Graphs?

- ▶ If graph is sparse, we can do better with Dijkstra from every vertex, provided weights are non-negative
 - ▶ Dijkstra: $O(|V||E| + |V|^2 \log(|V|))$ vs Floyd-Warshall $O(|V|^3)$
- ▶ Can we apply Dijkstra when graph has negative weight edges, but not negative cycles?
 - ▶ Graph reweighing

Graph Reweighting

Theorem. Given a function $h : V \rightarrow \mathbb{R}$, **reweigh** each edge $(u,v) \in E$ by $w_h(u,v) = w(u,v) + h(u) - h(v)$.

Then, for any two vertices, all paths between them are reweighed by the same amount
==> shortest path in reweighed graph is also shortest in original path

Proof. Let $p = v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k$ be a path in the original graph G . We have

$$w(p) = w(v_1, v_2) + w(v_2, v_3) + \dots + w(v_{k-1}, v_k)$$

On the reweighed graph G_h , we have

$$\begin{aligned} w_h(p) &= w_h(v_1, v_2) + w_h(v_2, v_3) + \dots + w_h(v_{k-1}, v_k) \\ &= w(v_1, v_2) + h(v_1) - h(v_2) + w(v_2, v_3) + h(v_2) - h(v_3) + \dots + w(v_{k-1}, v_k) + h(v_{k-1}) - h(v_k) \\ &= w(p) + h(v_1) - h(v_k) \end{aligned}$$

Same amount for every path that starts at v_1 , ends at v_k !

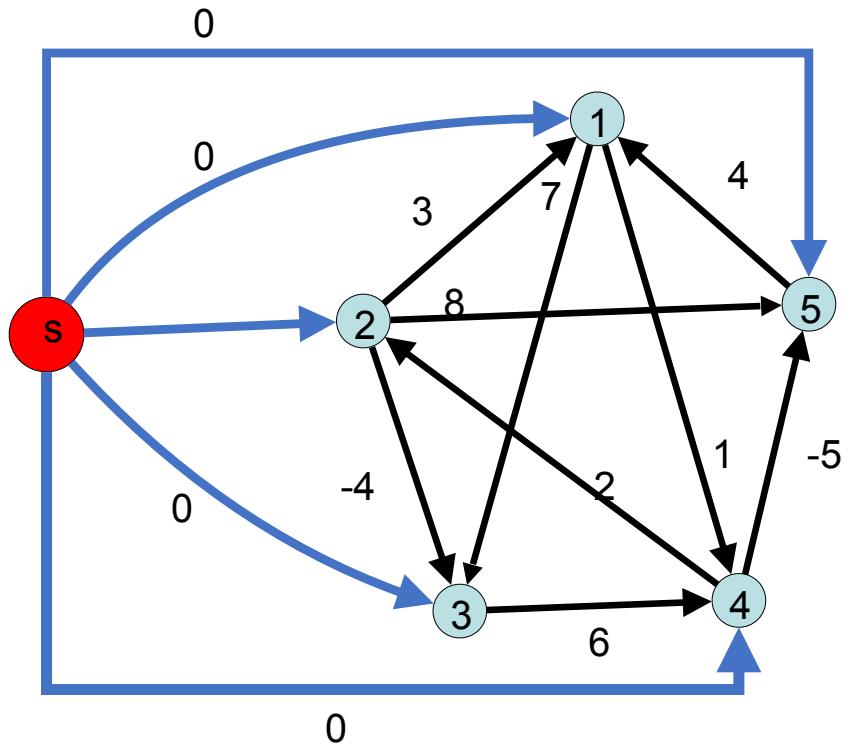
Can We Do Better for Sparse Graphs?

- So, want to find a reweighing function such that $w_h(u, v) \geq 0$ for all edges in E
 - Then, we could solve using Dijkstra's algorithm starting from each vertex
- **Johnson's algorithm**
 - Use an auxiliary shortest path problem solved with Bellman-Ford to use Bellman-Ford to find $h(v)$ such that $h(v) - h(u) \leq w(u, v)$, or determine that a negative-weight cycle exists.
 - Time = $O(|V| |E|)$
 - Resulting problem with $w_h(u, v) = w(u, v) + h(u) - h(v) \geq 0$

Johnson's algorithm

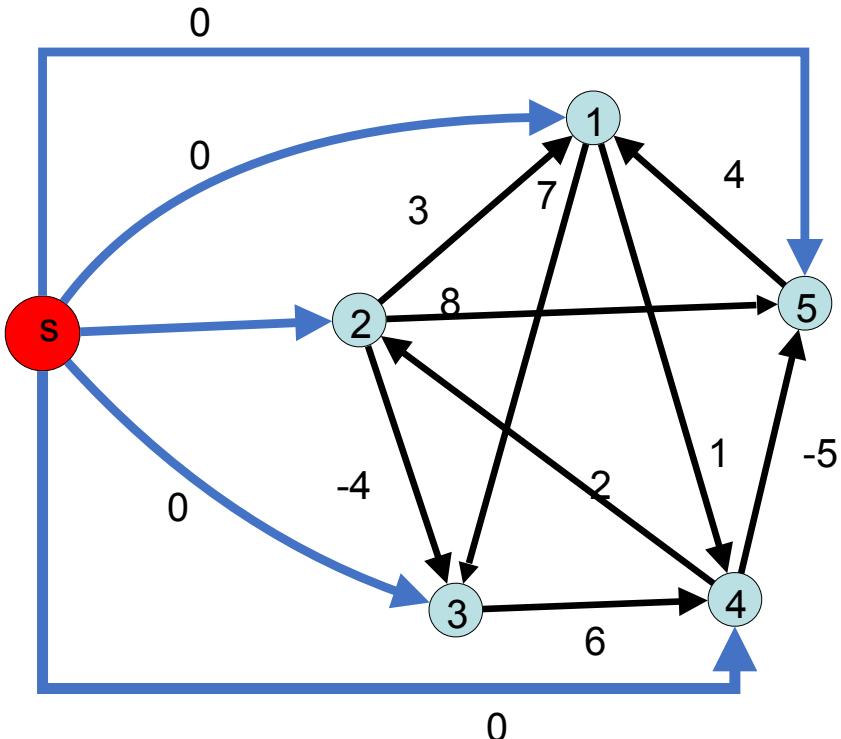
```
Johnson(G)
{   compute  $G'$ , where  $V[G'] = V[G] \cup \{s\}$  and
     $E[G'] = E[G] \cup \{(s, v) : v \in V[G]\}$ 
    if Bellman-Ford( $G'$ , w, s) = False
        then print "There is a neg-weight cycle"
    else for each vertex  $v \in V[G']$ 
        set  $h(v) = \delta(s, v)$  computed by Bellman-Ford algo.
        for each edge  $(u, v) \in E[G']$ :
             $w'(u, v) = w(u, v) + h(u) - h(v)$ 
        for each vertex  $u \in V[G']$ 
            run Dijkstra( $G, w', u$ ) to compute  $\delta'(u, v)$ 
        for each  $v \in V[G]$ 
             $d_{uv} = \delta'(u, v) - h(u) + h(v)$ 
return D
}
```

Johnson's algorithm



Main idea: Add a dummy vertex with a directed edge of weight 0 to all other vertices

Johnson's algorithm

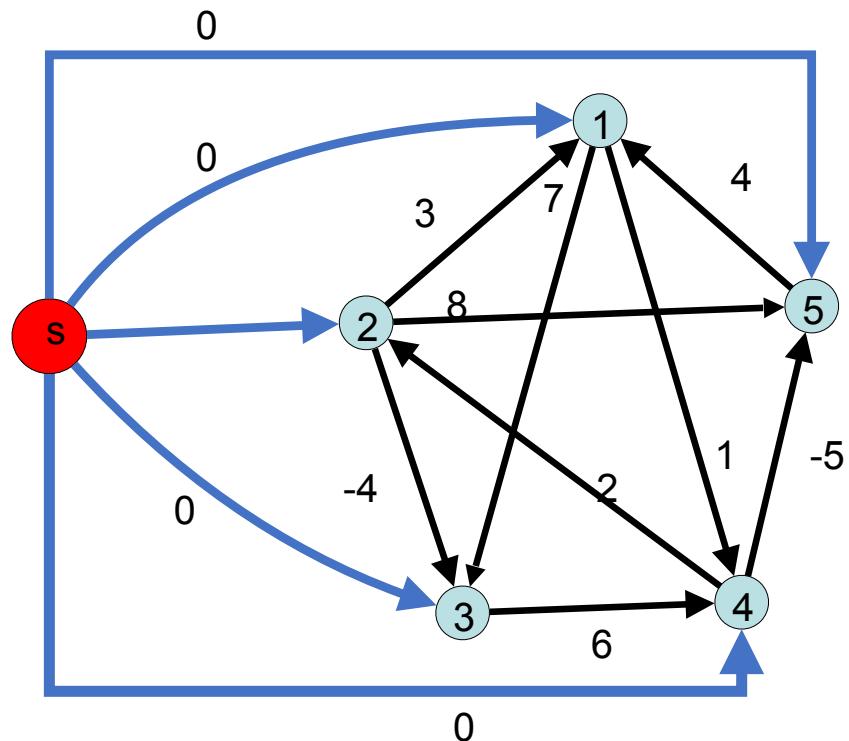


Use Bellman-Ford from s to get shortest distance to all other nodes:

Must be negative

	D[]	pred[]	Q[]
1	0	S	1
2	0	S	2
3	0	S	3
4	0	S	4
5	0	S	5

Johnson's algorithm



All-Pairs Shortest Paths

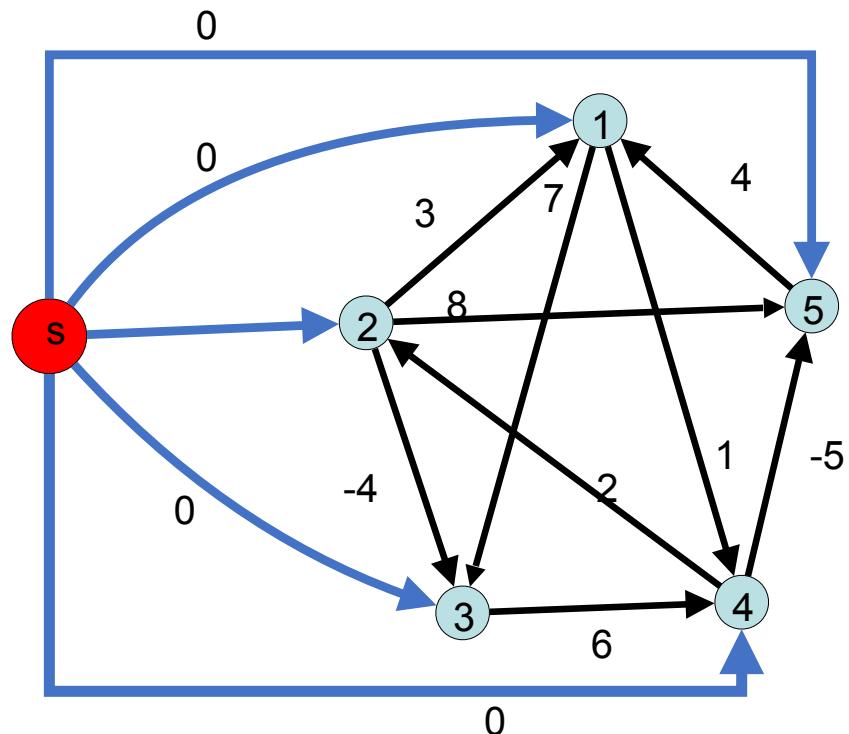
Pop 1

	D[]	pred[]	Q []
1	0	S	
2	0	S	1
3	0	S	2
4	0	S	3
5	0	S	4

Pop 2

	D[]	pred[]	Q []
1	0	S	
2	0	S	
3	-4	2	1
4	0	S	2
5	0	S	3

Johnson's algorithm



All-Pairs Shortest Paths

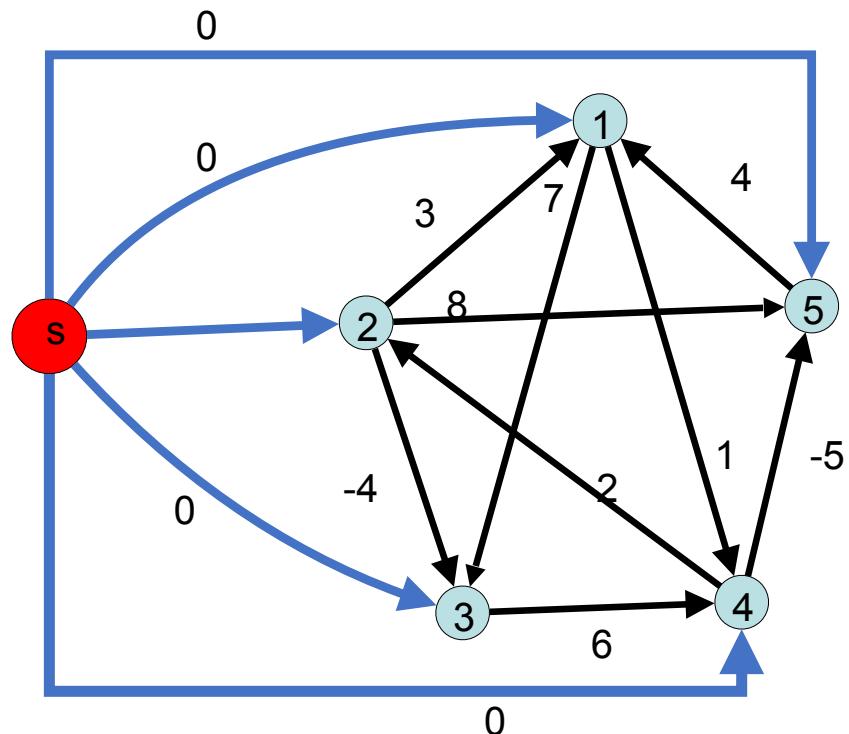
	D[]	pred[]	Q []
1	0	S	
2	0	S	
3	-4	2	
4	0	S	1
5	0	S	2

Pop 3

	D[]	pred[]	Q []
1	0	S	
2	0	S	
3	-4	2	
4	0	S	
5	-5	4	1

Pop 4

Johnson's algorithm



All-Pairs Shortest Paths

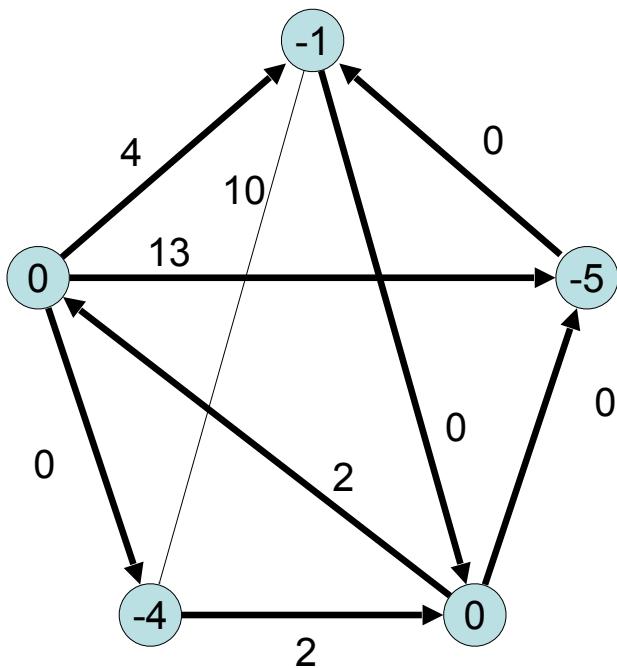
Pop 5

	D[]	pred[]	Q []
1	-1	5	1
2	0	S	
3	-4	2	
4	0	S	
5	-5	4	

Pop 1

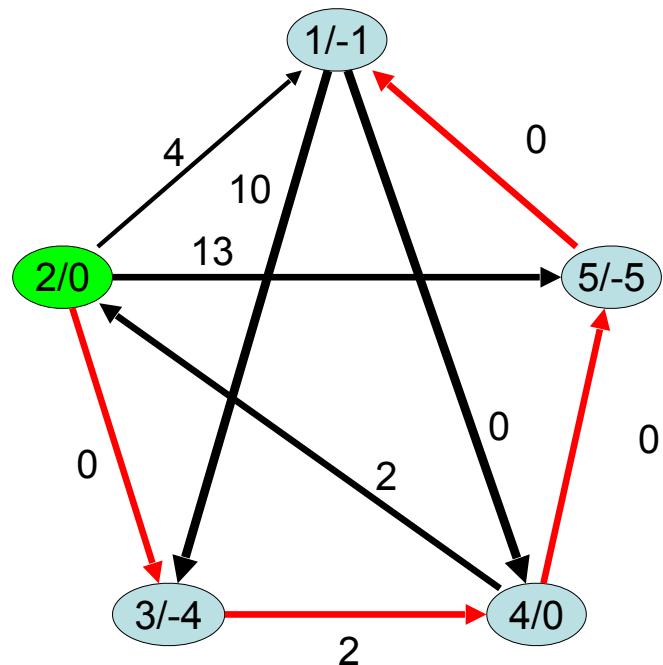
	D[]	pred[]	Q []
1	-1	5	
2	0	S	
3	-4	2	
4	0	S	
5	-5	4	

Johnson's algorithm



Use Distance from
previous problem
for reweighing

Johnson's algorithm



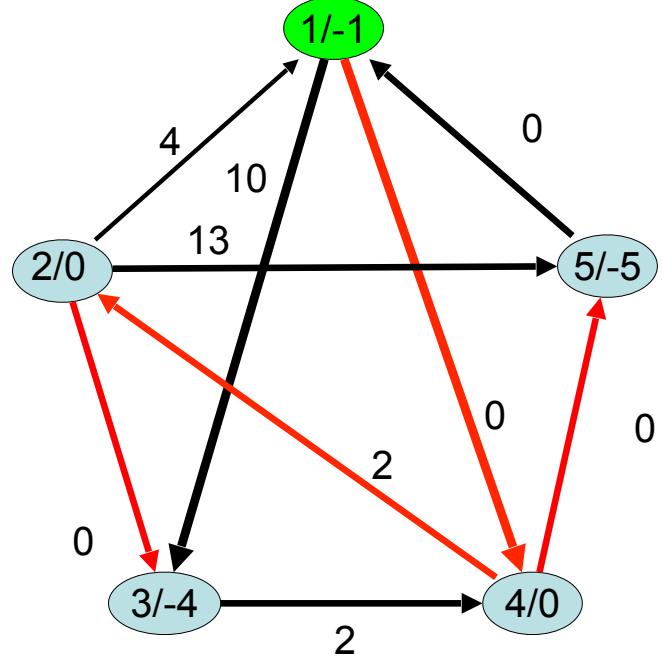
	D[]	pred[]	Q []
1	Inf	Null	
2	0	null	0
3	Inf	Null	
4	Inf	Null	
5	inf	Null	

	D[]	pred[]	Q []
1	Inf	Null	
2	0	null	
3	0	2	0
4	2	3	0
5	13	2	1

	D[]	pred[]	Q []
1	Inf	Null	
2	0	null	
3	0	2	
4	2	3	
5	2	4	1

	D[]	pred[]	Q []
1	Inf	Null	
2	0	null	
3	0	2	0
4	Inf	Null	
5	13	2	1

	D[]	pred[]	Q []
1	Inf	Null	
2	0	null	
3	0	2	
4	2	3	
5	2	4	1



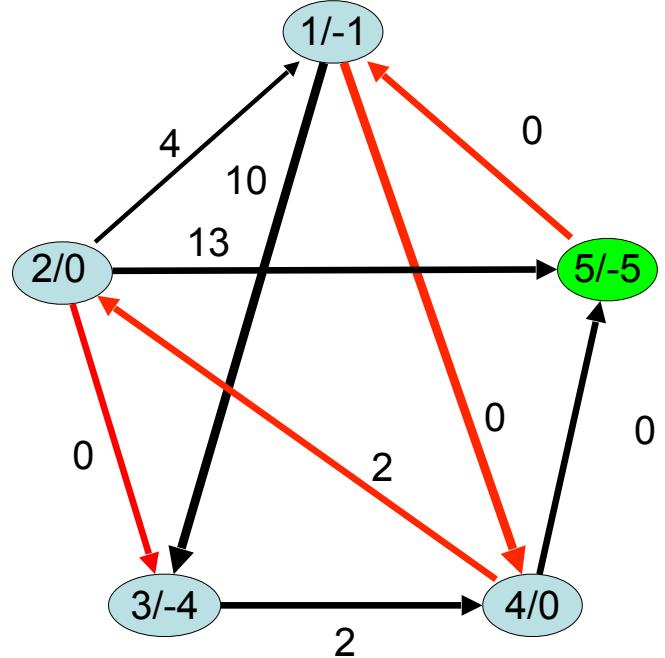
	D[]	pred[]	Q []
1	0	Null	0
2	Inf	null	
3	Inf	Null	
4	Inf	Null	
5	inf	Null	

	D[]	pred[]	Q []
1	0	Null	
2	2	null	2
3	10	1	1
4	0	1	
5	0	5	0

	D[]	pred[]	Q []
1	0	5	
2	2	null	0
3	0	2	0
4	0	1	
5	0	5	

	D[]	pred[]	Q []
1	0	Null	
2	Inf	null	
3	10	1	1
4	0	1	0
5	Inf	null	

	D[]	pred[]	Q []
1	0	Null	
2	2	null	0
3	10	1	1
4	0	1	
5	0	5	



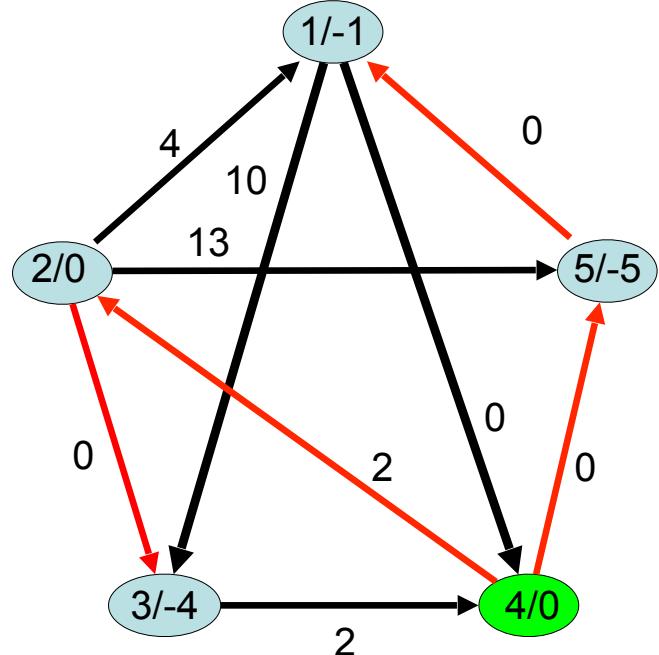
	D[]	pred[]	Q []
1	Inf	Null	
2	Inf	null	
3	Inf	Null	
4	Inf	Null	
5	0	Null	0

	D[]	pred[]	Q []
1	0	5	
2	Inf	null	
3	10	1	1
4	0	1	0
5	0	null	

	D[]	pred[]	Q []
1	0	5	
2	2	4	0
3	10	1	1
4	0	1	
5	0	null	

	D[]	pred[]	Q []
1	0	5	0
2	Inf	null	
3	inf	Null	
4	inf	Null	
5	0	null	

	D[]	pred[]	Q []
1	0	5	
2	2	4	0
3	10	1	1
4	0	1	
5	0	null	



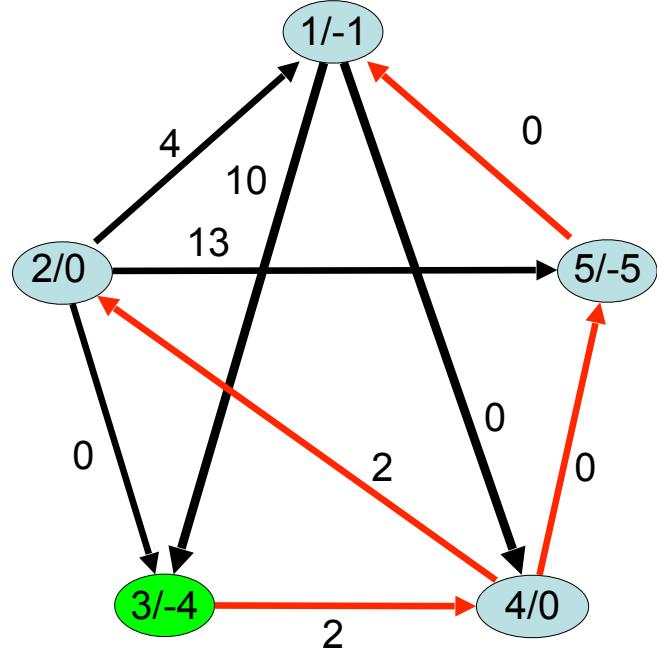
	D[]	pred[]	Q []
1	Inf	Null	
2	Inf	null	
3	Inf	Null	
4	0	Null	0
5	Inf	Null	

	D[]	pred[]	Q []
1	0	5	0
2	2	4	1
3	Inf	Null	
4	0	Null	
5	0	4	

	D[]	pred[]	Q []
1	0	5	
2	2	4	
3	2	3	0
4	0	Null	
5	0	4	

	D[]	pred[]	Q []
1	Inf	Null	
2	2	4	1
3	Inf	Null	
4	0	Null	
5	0	4	0

	D[]	pred[]	Q []
1	0	5	
2	2	4	0
3	10	1	1
4	0	Null	
5	0	4	



	D[]	pred[]	Q []
1	Inf	Null	
2	Inf	null	
3	0	Null	0
4	inf	Null	
5	Inf	Null	

	D[]	pred[]	Q []
1	Inf	Null	
2	4	4	1
3	0	Null	
4	2	3	
5	2	4	0

	D[]	pred[]	Q []
1	2	5	
2	4	4	1
3	0	Null	
4	2	3	
5	2	4	

	D[]	pred[]	Q []
1	Inf	Null	
2	Inf	null	
3	0	Null	
4	2	3	0
5	Inf	Null	

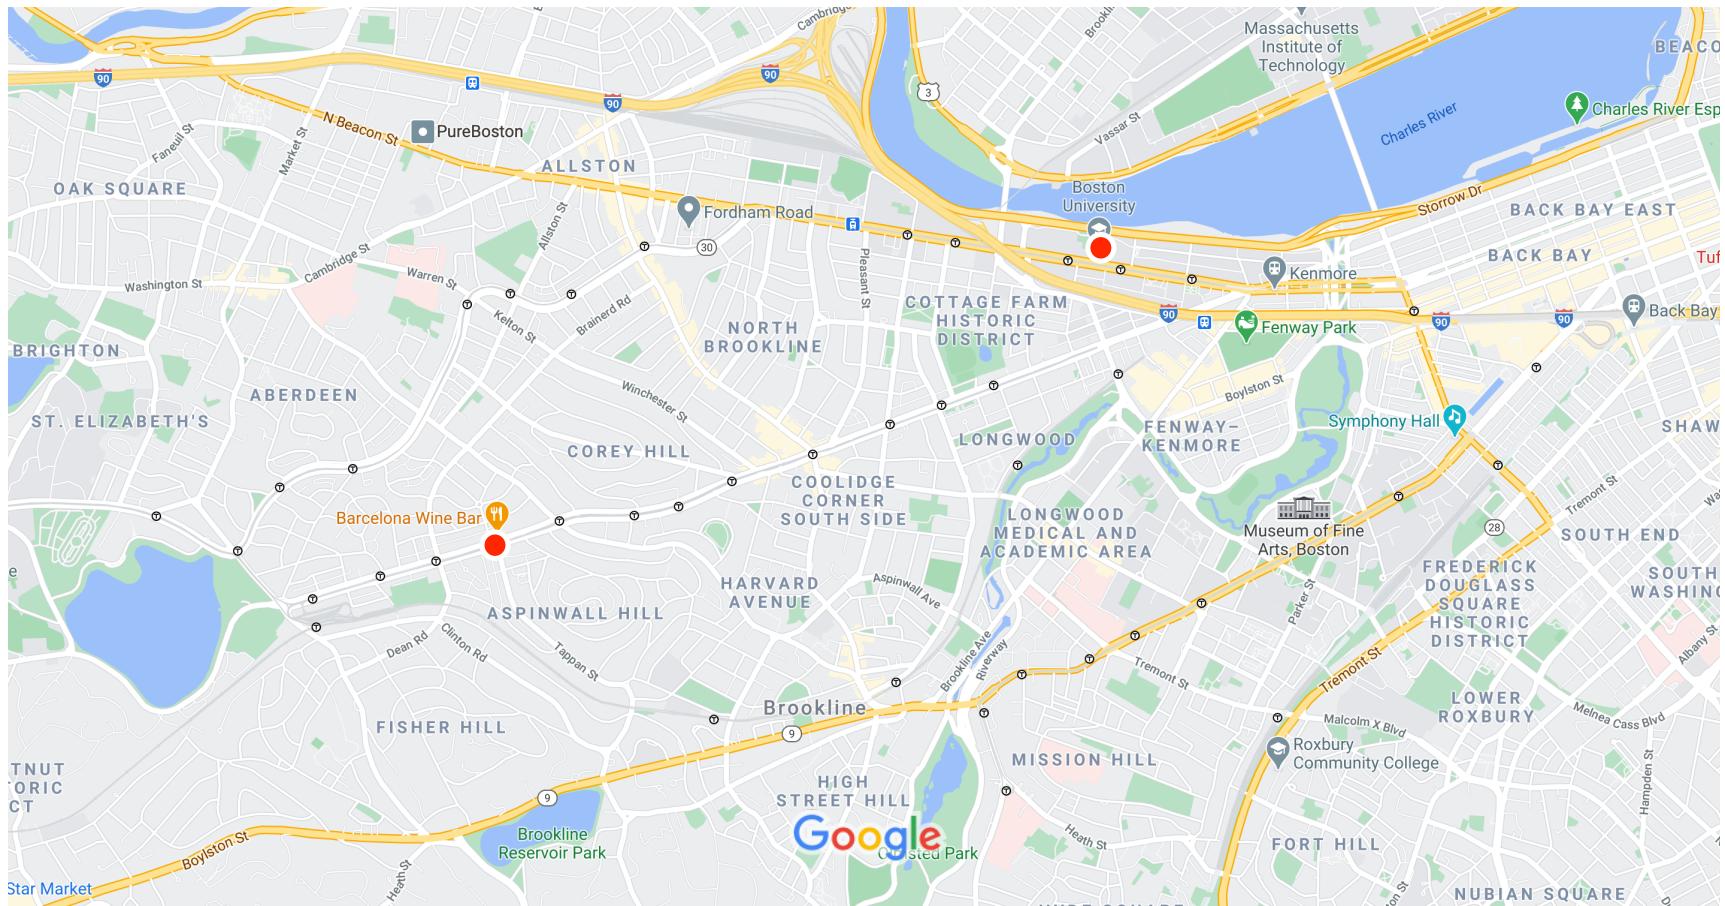
	D[]	pred[]	Q []
1	2	5	0
2	4	4	1
3	0	Null	
4	2	3	
5	2	4	

Problem: Shortest Path from s to a single t

- We only want a path to a destination
 - But all algorithms compute paths from s to every destination
 - When can we stop?
 - Can we focus on getting only to where we want to go?
 - This is the Google maps or Waymo problem
- New algorithm: A^* search
 - Will search in direction of t for shortest path
 - A modification of Dijkstra to bias the search
 - Guaranteed optimal under certain conditions

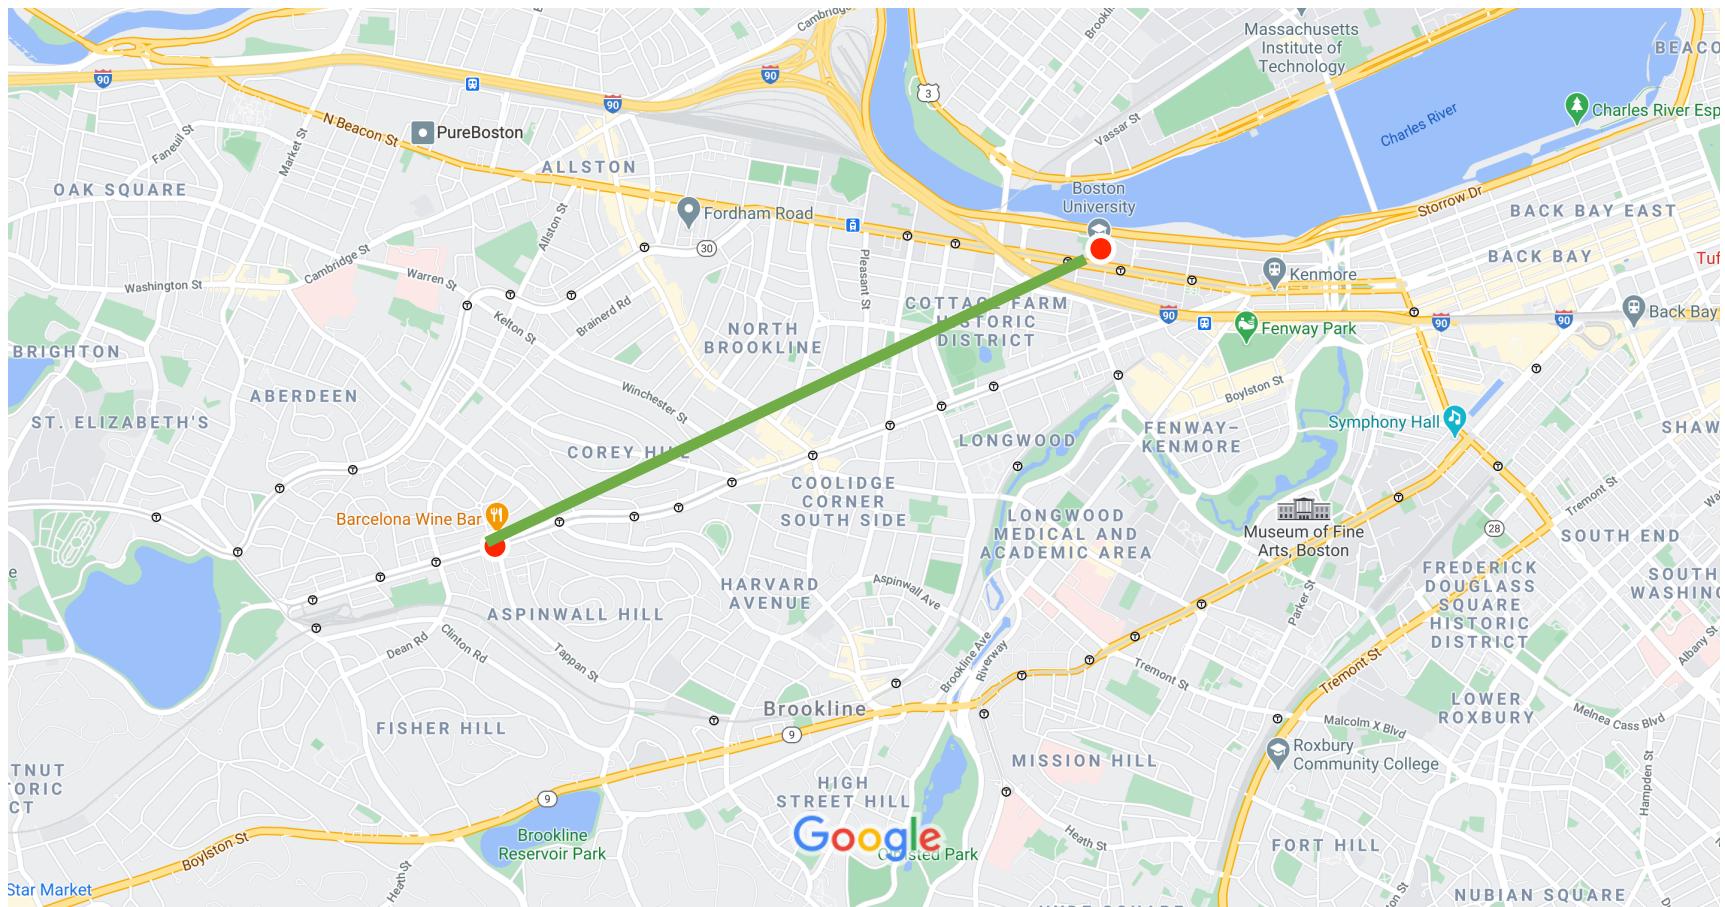
A* search

Example



A* search

Distance lower bound: Straight Path



Heuristic Distance Estimates

- Let $h : V \rightarrow \mathbb{R}$ be a heuristic function that approximates the distance from vertex v to target t
- We call h **admissible** if: for each vertex v : $h(v) \leq D(v,t)$, where $D(v,t)$ is the true shortest path distance
- We call h **consistent** if: For each (u,v) in E : $h(u) \leq w(u,v) + h(v)$.
 - For example, Euclidean distance to target on the plane
 - Euclidean distance is admissible and consistent

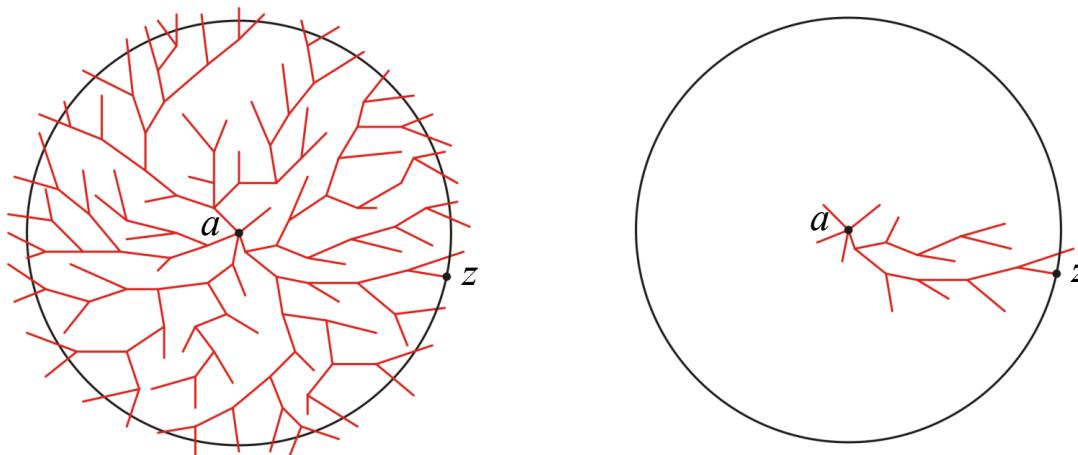
A* Search

- Simple idea: use heuristic $h()$ to reweigh edges
 - $w_h(u, v) = w(u, v) - h(u) + h(v)$
 - Measures deviation from ideal distance estimate by h
 - Consistent heuristics guarantee non-negative weights
- Solve using Dijkstra's algorithm
- Result: If heuristic is consistent and admissible, we get an optimal shortest path
 - Dijkstra's distance is biased towards vertices that have shorter distance estimates

Comparison with Dijkstra's Algorithm

Graphically, we can suggest the behavior of the two algorithms as follows:

- Suppose we are moving from a to z :



Representative search patterns for Dijkstra's and the A* search algorithms

A* search

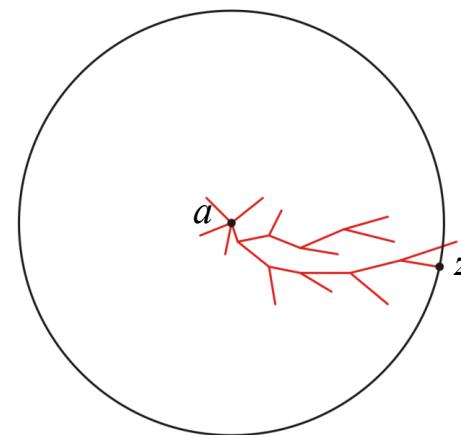
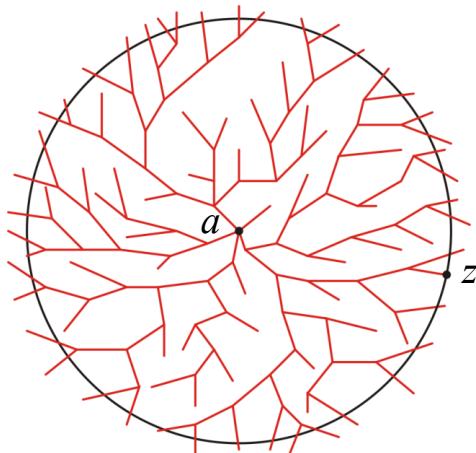
Comparison with Dijkstra's Algorithm

Dijkstra's algorithm is the A* search algorithm when

using the discrete distance

$$h(u, v) = \begin{cases} 0 & u = v \\ 1 & u \neq v \end{cases}$$

- No vertex is better than any other vertex



Representative search patterns for Dijkstra's and the A* search algorithms