

EC504 ALGORITHMS AND DATA STRUCTURES
FALL 2020 MONDAY & WEDNESDAY
2:30 PM - 4:15 PM

Prof: David Castañón, dac@bu.edu

GTF: Mert Toslali, toslali@bu.edu

Haoyang Wang: haoyangw@bu.edu

Christopher Liao: cliao25@bu.edu

New Problem: Large Databases

Organization	Database Size
WDCC	6,000 TBs
NERSC	2,800 TBs
AT&T	323 TBs
Google	33 trillion rows (91 million insertions per day)
Sprint	3 trillion rows (100 million insertions per day)
ChoicePoint	250 TBs
Yahoo!	100 TBs
YouTube	45 TBs
Amazon	42 TBs
Library of Congress	20 TBs

Source: www.businessintelligencelowdown.com, 2007.


- o How are these stored for efficient search? Oracle, SQL? Not in RAM...

Large Databases

- Use a BST?
 - Google – 33 trillion items, indexed by IP
 - Access time
 - Height = $\log(33 \times 10^{12}) = 44.9$
 - Assume 120 disk accesses per second: Search takes 0.37 seconds
- What does Oracle, SQL do?
 - Use better search trees to reduce disk access
- How about other solutions (e.g. hash tables, like Python dictionaries?)
 - Many data base queries imply sorting (hard in hash tables)
 - Huge data bases need to manage disk I/O (not fit in memory!)

Want to Exploit Memory Hierarchy

- The lower you go in the hierarchy, the less you want to go back
- Limit accesses to very small numbers



Idea: Multiway Search Trees


- Idea: allow a node in a tree to have many children
- Less disk access = less tree height = more branching
- As branching increases, the depth decreases
- An M-ary tree allows M-way branching
- Each internal node has at most M children
 - A complete M-ary tree has height that is roughly $\log_M(N)$ instead of $\log_2(N)$
 - If $M = 20$, then $\log_{20}(2^{20}) < 5$
- We can speedup the search significantly

Example

- Standard disk page size 8192 bytes
- Assume keys use 32 bytes, pointers use 4 bytes
- Keys uniquely identify data elements
- $32*(M-1)+4*M = 8192$
- $M = 228$ nodes in a page
- $\text{Log}_{228} 33 \times 10^{12} = 5.7$ (disk accesses)
- Each search takes 0.047 seconds

B-Tree (actually, B+-Tree)

- A B+-tree of order M is an M-ary tree with the following properties
 - Data items are stored at the leaves
 - Non-leaf nodes store up to M-1 keys; keys at node are sorted
 - Non-leaf nodes have between $\lceil \frac{M}{2} \rceil$ and M links to children, except for root
 - The root is either a leaf or has from two to M children, 1 and M-1 keys
 - All leaves at the same depth, and contain the data items for the tree
 - Leaves have between $\lceil \frac{L}{2} \rceil$ and L data items
- Requiring nodes to be half full avoids degenerating into binary tree



B-Tree vs B⁺-Tree

- B Trees originally proposed in 1970 (Bayer and McCreight)
 - Data items are stored at interior nodes and at the leaves
 - Problem: data items can be much larger than simple navigation nodes, making this impractical!
 - Makes interior nodes larger, less keys possible, increased height
- B⁺ trees store all data at leaves
 - Leaf nodes store less data items (L vs M-1) because data items are larger
 - Shorter trees, better fit to hierarchical memory
 - To make things fast, first couple of levels of B⁺ trees kept in main memory
 - Choosing L:
 - Assuming a data element requires 256 bytes; leaf node 8192 bytes implies L=32
 - Each leaf node has between 16 and 32 data elements

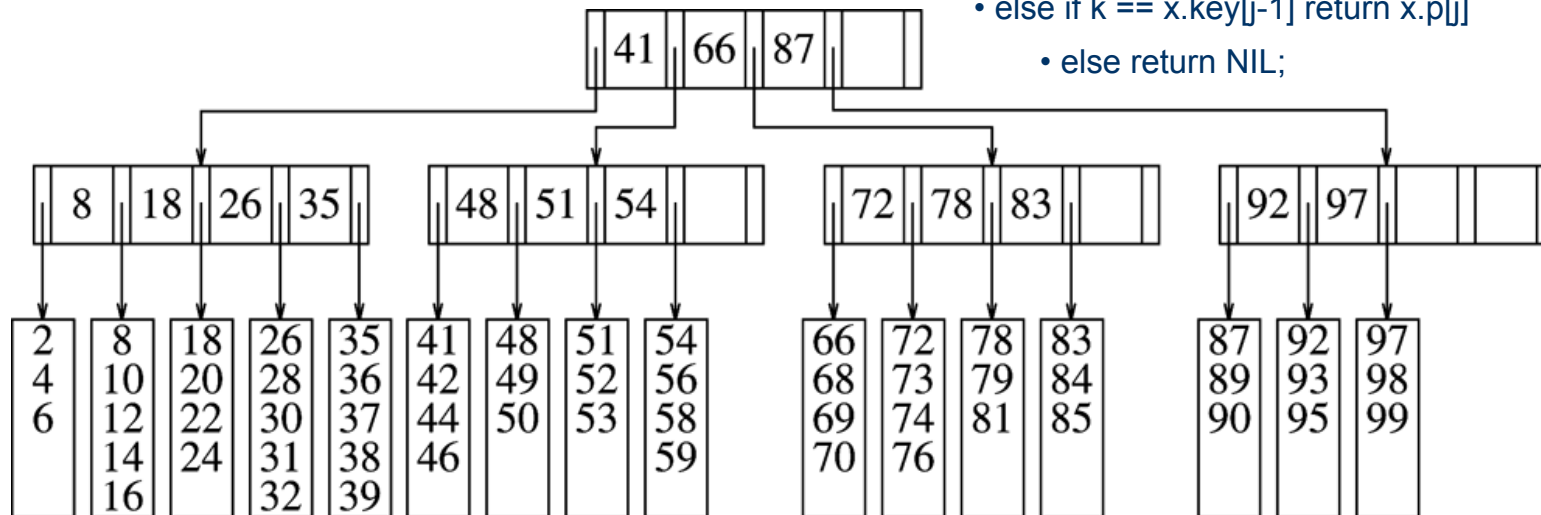
B+ Tree Search

- B+ tree of order 5
 - Nodes have 2-4 keys and 3-5 children, leaves have 3-5 data elements
- Example: Search for 53, search for 41
 - Searching node can be done with binary search
 - Some ambiguity, resolved by convention: pointer to the right of 41 includes keys $41 \leq k < 66$

- Node: $x.n$ = number of keys. $x.key[j]$ = j-th key
 - $x.c[j]$ = pointer to j-th child node; $x.leaf$ = Boolean,
 - $x.p[j]$ = pointer to the record corresponding to $x.key[j]$ in leaf

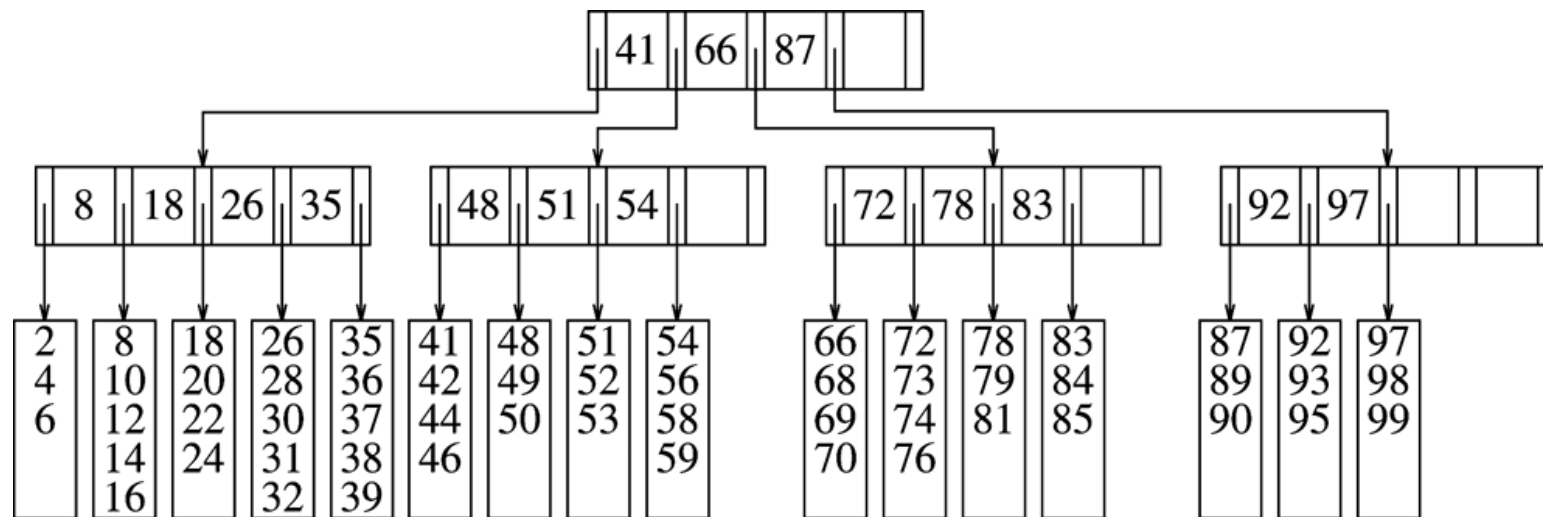
Search(node, k) for k starting at node x:

- $j = 0$; while $j < x.n$ and $k \geq x.key[j]$: $j++$;
- if not $x.leaf$: Search($x.c[j]$, k)
- else if $k == x.key[j-1]$ return $x.p[j]$
- else return NIL;



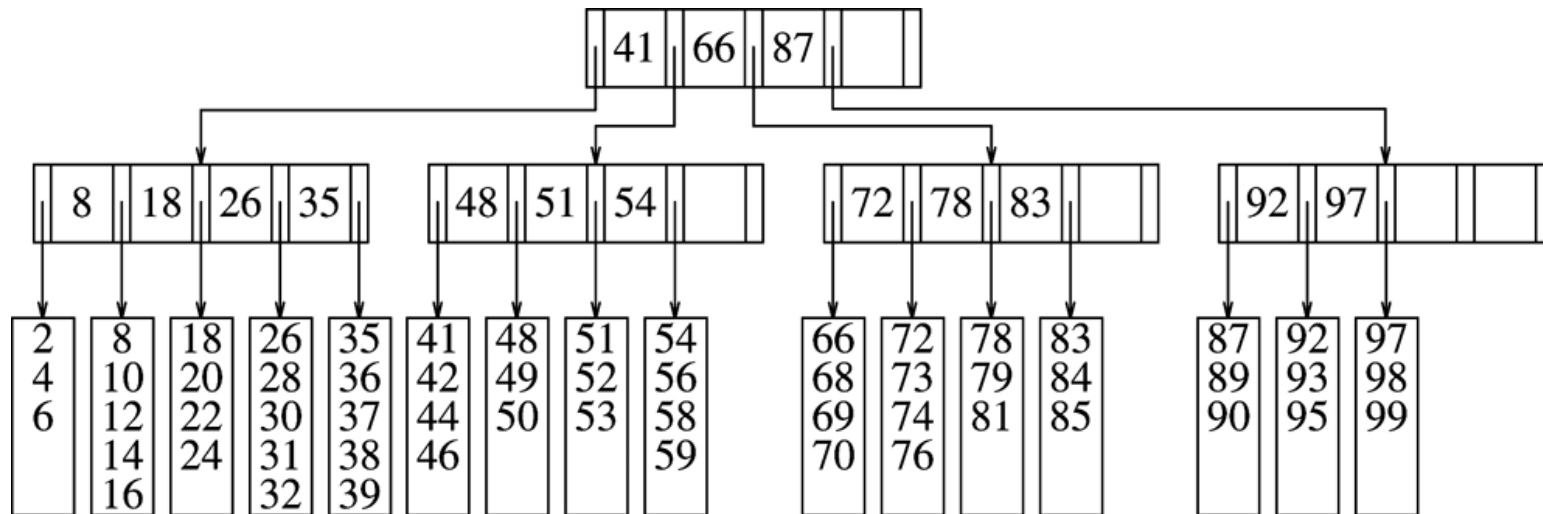
B+ Tree Insert

- Insert k
 - Navigate to leaf where key should be
 - **Case 1:** if there is room in that leaf (less than L keys): just add it there, sort key with other keys
 - e.g. Insert 55




B⁺ Tree Insert - 2


- Insert k
 - Navigate to leaf where key should be
 - **Case 2:** Leaf has L keys already (e.g. insert 40)
 - Need to split the leaf; split the leaf, promote middle key to parent node
 - May need to split parent if no room there...



B+ Tree Insert: Splitting


- Insert 40: split leaf, leave keys in tree. Move rightmost key in right split up to parent for navigation, add navigation link in parent
- Splitting interior node: remove move middle key of saturated node and add to parent, break node into two nodes, add navigation links in parent





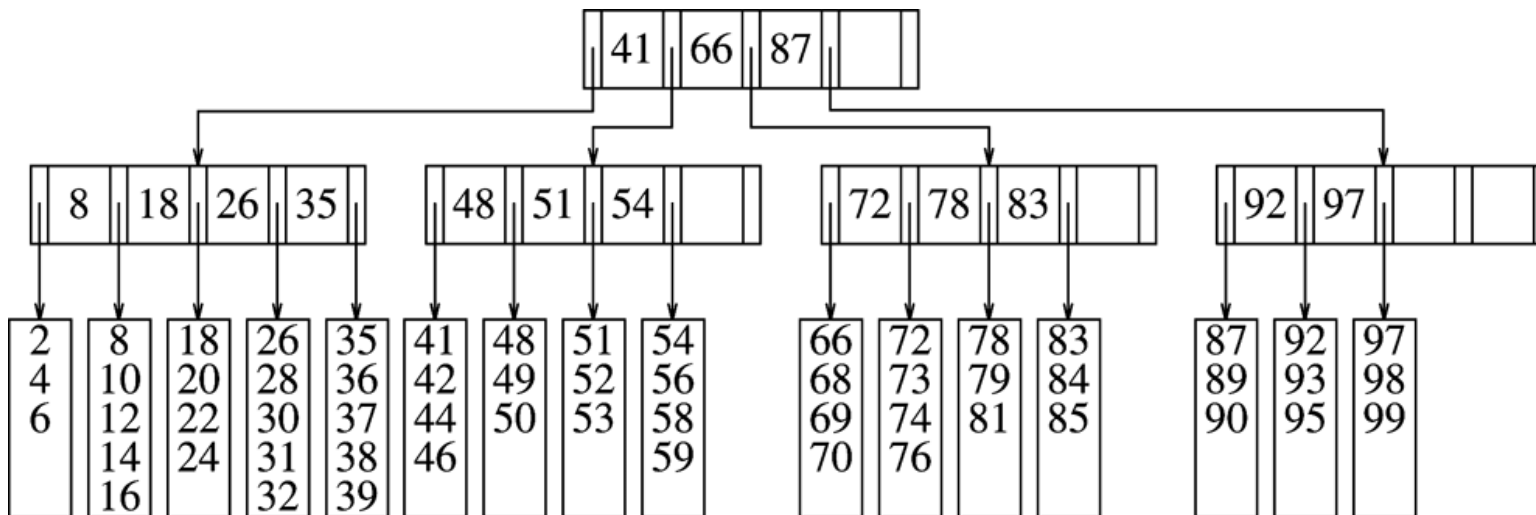
B+ Tree Delete

- Delete k
 - Navigate to leaf where key should be
 - **Case 1:** leaf has more than minimum number: Just delete it!
 - e.g. Delete 31



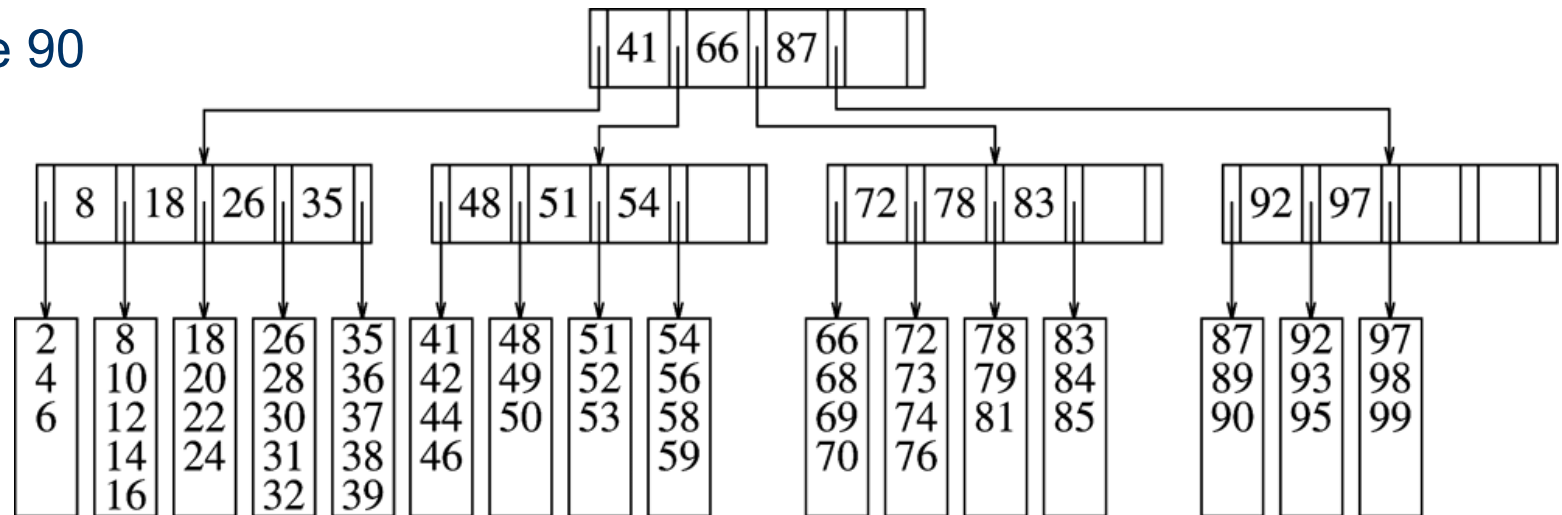
B+ Tree Delete

- Delete k
 - Navigate to leaf where key should be
 - **Case 2:** leaf has minimum number: Delete key. If one of immediate neighbor siblings has more than minimum number, transfer key to leaf, adjust value of key in parent
 - e.g. Delete 53



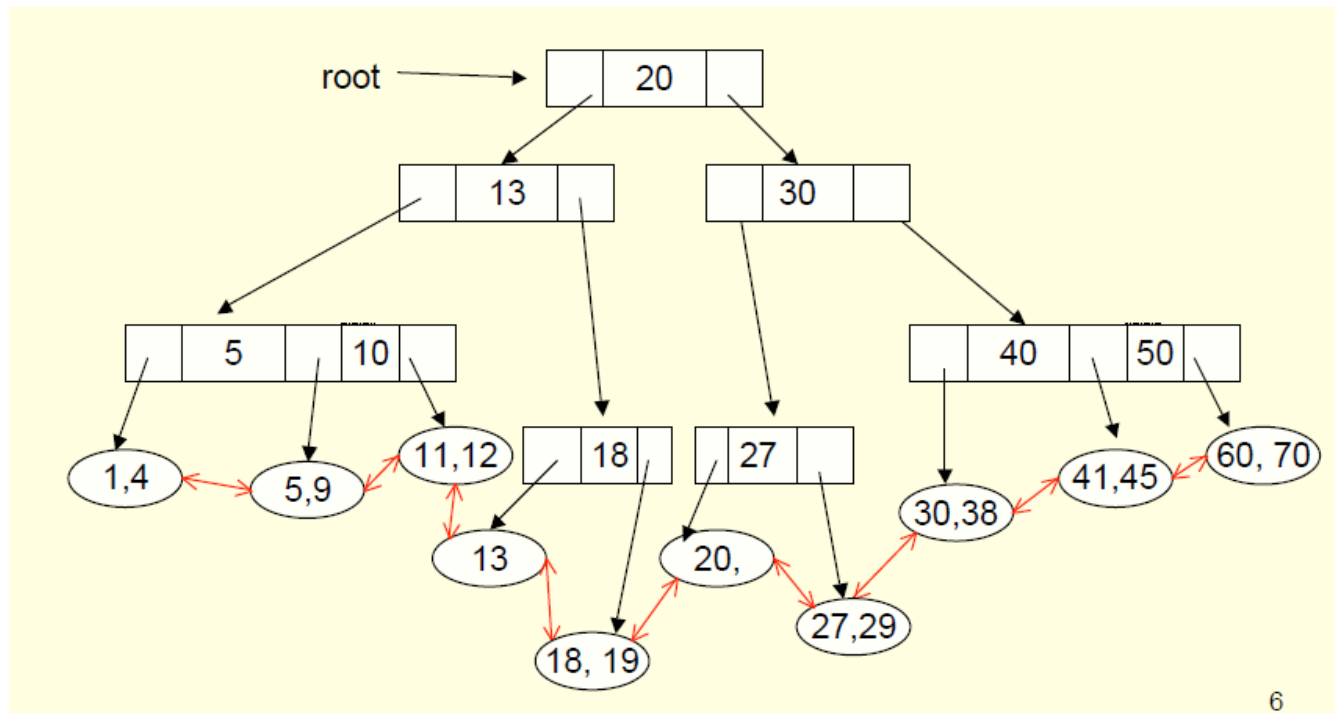
B+ Tree Delete

- Delete k
 - Navigate to leaf where key should be
 - **Case 3:** leaf has minimum number: Delete key. If immediate neighbor siblings have only minimum number, merge with one of neighbors, remove key in parent (May cause underflow in parent, so continue delete)
 - e.g. Delete 90



B+ Tree Delete

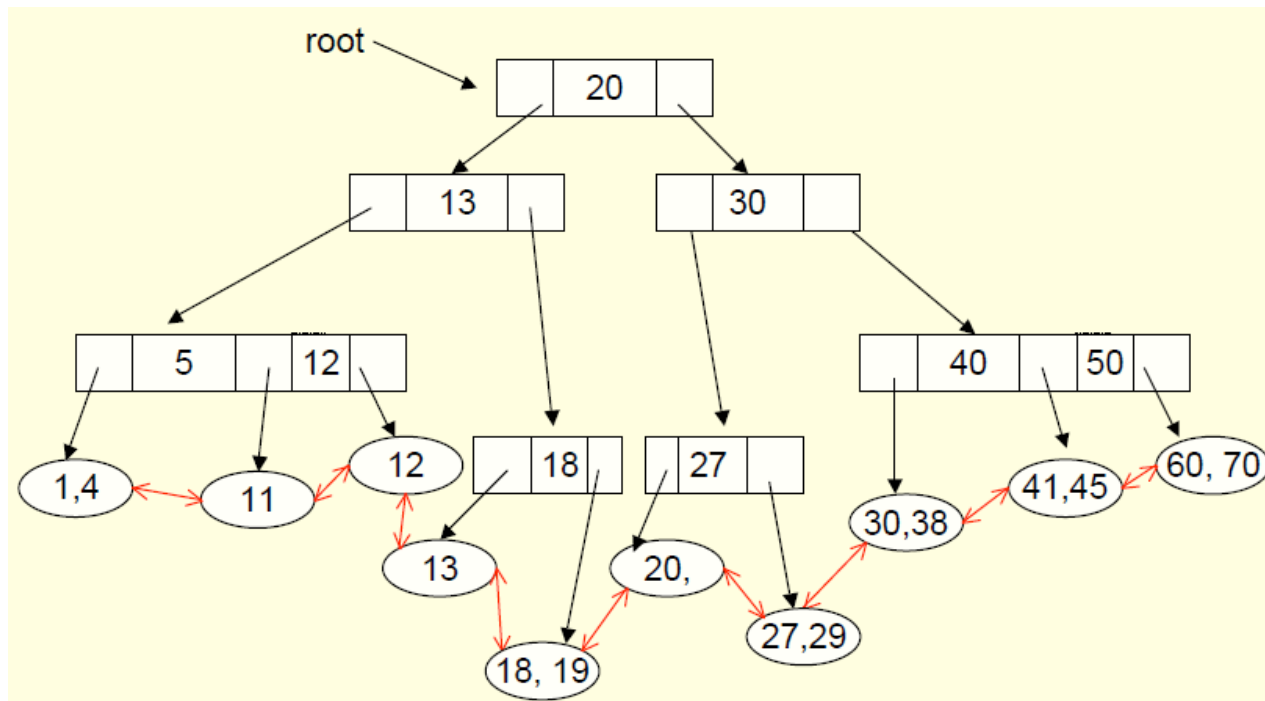
- Delete 5
 - No problem, leaf has extra key



6

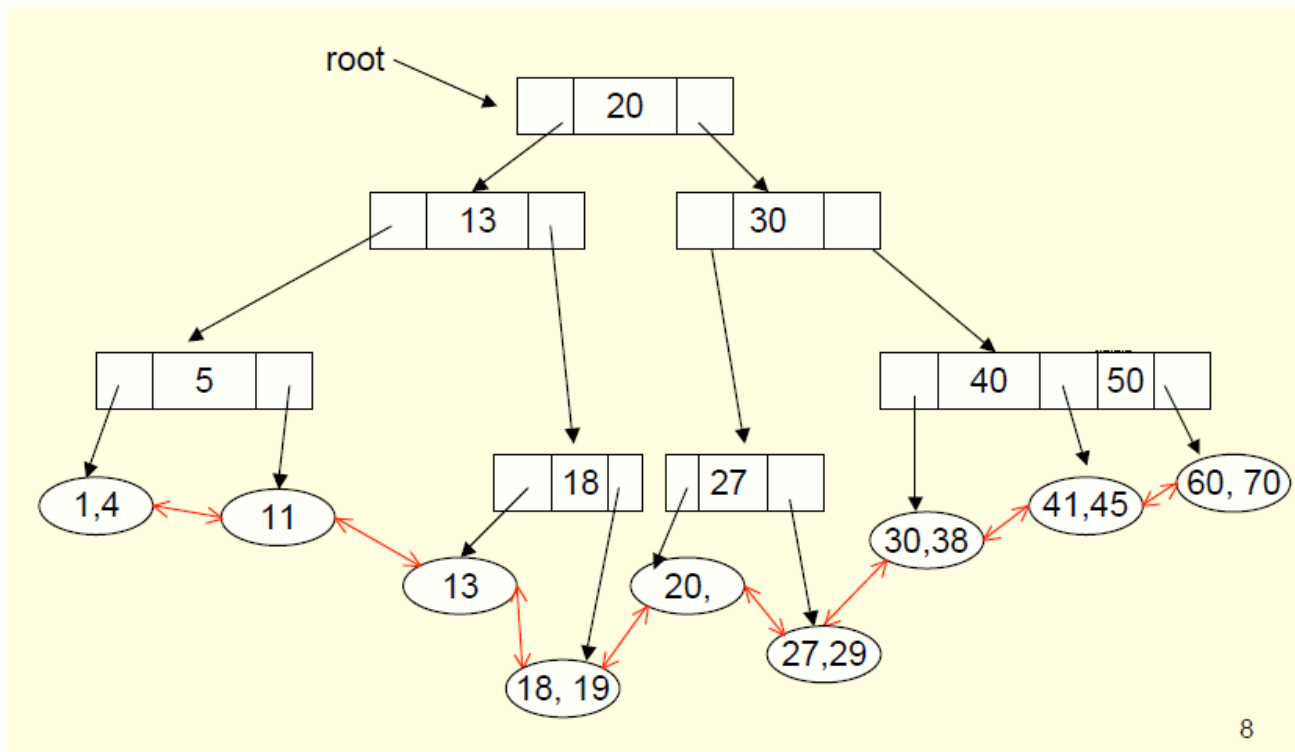
B+ Tree Delete

- Delete 5
 - No problem, leaf has extra key
- Delete 9: Now have underflow, must borrow key




B+ Tree Delete

- Delete 12
 - Underflow, must delete node



B+ Tree Delete


- Delete 4
 - No problem, extra key
- Delete 11
 - Big problem!
 - Underflow parent
 - Parent must merge



B+ Tree Delete


- Merge parent
 - Grandparent now has underflow

=>grandparent not full enough



B+ Tree Delete

- Merge grandparent
 - Root has underflow, delete root



11

B⁺ Tree Summary

- Insert, delete, find are $O(\log(n))$
 - Navigate to leaf where key should be
- Often, use pointers between neighbor leaves to obtain sorted records
- Maintain sorted keys inside node for easier navigation
- Variations:
 - B* tree very similar to B+ tree
 - instead of splitting a page in 1/2 when it overflows, it gives some records to its neighbor siblings.
 - If neighbor is full, then 2 nodes split into 3. This makes nodes at least 2/3 full.
 - Similarly, in deletions, one starts to shift records when number is less than 2/3 of maximum

Priority Queues

- Queues are a standard mechanism for ordering tasks on a first-come, first-served basis
 - However, some tasks may be more important or timely than others (higher priority)
- Priority queues
 - Store tasks using a partial ordering based on priority
 - Ensure highest priority task at head of queue
- Heaps are the underlying data structure of priority queues

Priority Queue

- Main operations
 - **Insert** (i.e., enqueue)
 - **deleteMin** (i.e., dequeue)
 - Finds the minimum element in the queue, deletes it from the queue, and returns it
 - **getMin**: Find the highest priority job
 - **deleteAny**: Delete a job from the queue
 - **Create**: Create a priority queue from a list of jobs
 - **Merge**: merge two priority queues into one
 - **decreasePriority**: decrease priority of existing job in queue
- Desired Performance
 - Goal is for operations to be fast: $O(1)$ to $O(\log(n))$
 - Will be able to achieve $O(\log N)$ time insert/deleteMin


Simple Implementations

- Unordered list
 - $O(1)$ insert
 - $O(n)$ deleteMin
 - $O(n)$ for several others
- Ordered list
 - $O(N)$ insert
 - $O(1)$ deleteMin
 - Still $O(n)$ for others
- Balanced BST (red-black tree, splay tree)
 - $O(\log(n))$ insert and deleteMin
 - $O(\log(n))$ decrease priority
- Observation: We don't need to keep the priority queue completely ordered

Binary Heap

- A binary heap is a binary tree with two properties
 - Structure property
 - A binary heap is a complete binary tree
 - Each level is completely filled
 - Bottom level may be partially filled from left to right
 - Heap-order property
 - Parent node must have key less than or equal to the keys of its children
- Complete binary tree —> easy implementation as array
 - Height of a complete binary tree with N elements is **Floor**[logN]

Binary Heap Example




- Map nodes into array positions, starting from 0:
 - Root node: element 0
 - Children of node in position k : are in positions $2k+1, 2k+2$
 - Parent of node in position k : is in position $(k-1)/2$

Heap Insert

- Insert new element into the heap at the next available slot (“hole”)
 - According to maintaining a complete binary tree
- Then, “percolate” the element up the heap while heap-order property not satisfied
 - Upheap operation

Heap Insert

Insert 14:



Creating the hole and building the hole up

Upheap(k)

If (k=0) return;

If $V[k] < V[(k-1)/2]$

temp = $V[(k-1)/2]$

$V[(k-1)/2] = V[k]$

$V[k] = \text{temp};$

$k = (k-1)/2$


Upheap(k)

else return;

Heap DeleteMin

- Minimum element is always at the root
 - $O(1)$ to find
 - On deleteMin, heap decreases by 1 in size
- Operations
 - Remove element in array[0]
 - Move last element into array[0], shrink array size by 1
 - Percolate down while heap-order property not satisfied
 - Downheap(k)

Heap DeleteMin



downheap(k,n): n = # of elements i

$V = V[k]$

$j = (2*k+1)$ // left child

while ($j \leq n$) {

 if $j < n-1$ and $V[j] > V[j+1]$:

$j = j+1$

 if $V[k] > V[j]$:

$V[k] = V[j]$

$k = j$;

$j = 2*j + 1$ // leftchild

 else exit

$V[j] = V$;

Building a Heap from Array


- Construct heap from initial array of n items
 - Solution 1: Perform n inserts
 - $O(n)$ average case, but $O(n \log(n))$ worst-case (reverse sorted)
 - Solution 2
 - Perform downheap for all nodes from position $n/2 - 1$ to 0.
 - Buildheap(n): for ($i = n/2 - 1$; $i \geq 0$; $i--$) downheap(i, n)
 - $O(N)$ worst case!

Complexity: $1/2$ elements don't move, $1/4$ move 1, $1/8$ move 2, ...

$$T(n) = \frac{n}{4} \left(1 + 2 * \frac{1}{2} + 3 * \frac{1}{4} + \dots \right) = \frac{n}{4} \sum_{k=1}^{\log(n)} \frac{k}{2^{k-1}} \leq n \Rightarrow \in \Theta(n)$$


Examples

Insert 43, then 18



Examples

Building a Heap from an array



Complexity of Heap Operations

- Height of heap: $O(\log(n))$
- insert: $O(\log(n))$
 - 2.607 comparisons on average, i.e., $O(1)$
- deleteMin: $O(\log(n))$
- decreaseKey: Find key, change value, upheap: $O(\log(n))$
- increaseKey: Find key, change value, downheap: $O(\log(n))$
- remove: Find key, swap last element into key position, downheap: $O(\log(n))$
- buildHeap: $O(n)$
- findMinimum: $O(1)$

- Merge two heaps: $O(n) \rightarrow$ link two arrays and buildHeap

Wishes: Merge of $O(1)$, decrease key of $O(1) \rightarrow$ useful in future algorithms...

Fast Sorting Algorithm: HeapSort

- Input: array of numbers (n)
- Step 1: Build heap: $O(n)$
- Step 2: for k in 1 to n , deleteMin(): Complexity $O(n \log(n))$
- Heapsort complexity: $O(n \log(n))$: as small as the fastest comparison based sort algorithms.
- Only drawbacks: Not stable, and average case is $O(n \log(n))$: not opportunistic

Priority Queue Applications

- Operating Systems: task scheduling
- Graph algorithms: Dijkstra's shortest path algorithm, Prim's Minimum Spanning Tree, A* search, Branch and Bound, ...
- Discrete-event simulation: Have fast access to next event
- Heapsort complexity: $O(n \log(n))$: as small as the fastest comparison based sort algorithms.
- Only drawbacks: Not stable, and average case is $O(n \log(n))$: not opportunistic